



THE STABILITY GAP AS A TRAJECTORY PROBLEM: A BENT DESCENT PATH, ITS OVERSHOOT, AND ONE GATE IN TWO DESIGNS

Bachelor's Thesis

Arthur Ernst Henrik Reuss, s5582253, a.e.h.reuss@student.rug.nl,
Supervisor: Gido van de Ven

Abstract: Replay-based continual learning is usually judged by the accuracy it reaches *after* each task. Continual evaluation exposes a failure this view hides: the *stability gap*, a sharp but transient drop in past-task accuracy at the start of every new task. The gap persists even when the optimiser descends the exact joint loss, so its cause lies in the trajectory, not the objective. This thesis explains the gap on two levels. A reference path exists along which the past-task loss rises monotonically to its unavoidable joint-optimum level; everything above that level is avoidable in principle. Steepest descent after an abrupt task switch leaves this path, bowing out of the old task's basin and returning from outside (level one, the trajectory problem), and real optimisers additionally overshoot whatever path they follow, pushed by an imbalanced first step and finite-buffer gradient noise (level two, the overshoot problem). Written in a unified update rule, replay methods that modify the update counter both problems with the same tool: a *gate* on the new-task gradient, with the replay gradient restoring at full strength. We build and compare the two pure gate designs. The λ -*curriculum*, a scalar feedback gate, ramps the new-task loss weight up from zero so the optimiser tracks the moving valley floor of the reference path; *asymmetric preconditioned experience replay*, a directional feedforward gate, filters the new-task gradient through past-task curvature, braking it only where the past task is steep. On rotated MNIST the two levels separate cleanly: removing the overshoot drivers leaves the arc, fixing the path leaves the overshoot, and doing both removes the gap while the permanent stability-plasticity cost remains. With momentum the curriculum cuts the gap to a few points of depth at vanilla-ER accuracy. The directional gate lowers the gap monotonically as its filter strengthens, at no plasticity cost but far higher compute; on an exact-gradient diagnostic its strongest setting fails late and abruptly, consistent with blind spots of its sampled curvature metric, and momentum re-inflates its per-step filter. Which gate wins is therefore set by the momentum regime: without momentum the directional gate is better on both stability and accuracy; with momentum, standard in practice, the scalar gate takes over. On CIFAR-10 with a ResNet-18 the curriculum cuts gap depth several-fold below vanilla experience replay at matched final accuracy and negligible compute cost, under two normalisation schemes and on corruptions it was never tuned on.

Contents

1	Introduction	3
2	Background and Related Work	4
2.1	Continual Learning and Experience Replay	4
2.2	The Stability Gap	4
2.3	One Update, Many Methods	4
2.4	Low-Loss Paths and Mode Connectivity	5
3	The Stability Gap as a Trajectory Problem	5
3.1	The Reference Path and the Tracking Bound	5
3.2	The Two-Level Gap: Bowing and Overshoot	6
3.3	Two Gates: Attenuating the New-Task Gradient	6

4	Research Questions	7
5	Methods	7
5.1	Experimental Setup	7
5.2	Experiments	8
6	Results	10
6.1	The Two Levels, Separated (RQ1)	11
6.2	The Scalar Feedback Gate (RQ2)	12
6.3	The Directional Feedforward Gate (RQ3)	14
6.4	Reading the Momentum Crosses Together	15
6.5	Generalisation to CIFAR-10 (RQ4)	16
6.6	Summary of Findings	17
7	Discussion	17
7.1	Why the gap needs two levels	17
7.2	Two gates, two ways to divide the work	18
7.3	What the account changes in the literature	19
7.4	Limitations	19
7.5	Outlook	20
	Statement on the Use of AI Tools	20
A	Implementation Details	23
A.1	Datasets	23
A.2	Models and Training	23
A.3	Experience replay	23
A.4	Asymmetric PER	24
B	Metric Definitions	24
C	Complete Results	25
C.1	rot-MNIST: Decomposition (D-series)	25
C.2	rot-MNIST: λ -Curriculum (C-series)	26
C.3	rot-MNIST: Buffer-Fidelity Diagnostics	27
C.4	rot-MNIST: Asymmetric PER (P-series)	27
C.5	rot-MNIST: Symmetric vs. Asymmetric Preconditioning	27
C.6	CIFAR-10: Headline Block	28
C.7	CIFAR-10: Generalisation Block	28
C.8	CIFAR-10: Momentum Cross and Compute Cost	30
C.9	Long-Sequence rot-MNIST (L-series)	31
D	Derivation of the Trajectory Bound	32

1 Introduction

Deep neural networks learn powerful representations from static, independent, and identically distributed data, but a vulnerability appears as soon as training becomes sequential: the network overwrites previously acquired knowledge to accommodate new data. This phenomenon, catastrophic forgetting [18, 21], is a major hurdle in machine learning. Retraining from scratch solves the issue, but it is too expensive for systems that need to adapt continuously. Continual learning (CL) studies how to learn from a sequence of tasks without that cost, and its core difficulty is the stability–plasticity dilemma: the model must stay plastic enough to learn new distributions while staying stable enough to preserve old ones [19].

Replay, the family this thesis studies, stores a small buffer of past examples and trains on them alongside the new data. Replay methods reduce long-term forgetting well, but until recently they were evaluated only at task boundaries, which hides what happens within a task. Lange et al. [16] evaluated the network after every update and identified the *stability gap*: a severe but transient collapse in past-task accuracy that occurs immediately when a new task begins, followed by a slow recovery. This matters for two reasons. A deployed system is queried throughout training, so a model caught mid-transition commits errors its final checkpoint never would. And because the model eventually recovers, past knowledge is disrupted rather than lost, which means the disruption itself is the target for repair.

The discovery raised a natural question: would a better approximation of the joint loss over all tasks remove the gap? Hess et al. [11] showed that the gap persists even under incremental joint training, where the model has access to all data from all tasks at every step. A perfect objective still produces transient forgetting. The problem therefore lies in *how* the optimisation proceeds, and this thesis locates it on two levels: the descent path itself is bent, and real optimisers overshoot whatever path they follow.

Two levels. Figure 1.1 visualises the transition as a switch from a learned task (Task A) to a new task (Task B), a conceptual reduction of a high-dimensional surface using the loss-landscape lens of Kao et al. [14]. *Level one* is geometric: steepest descent on the abruptly assembled joint loss (green curve) bows out of Task A’s basin and returns from outside [14], so the Task-A loss rises above its final level and comes back, its recovery slow and lingering, the *hanging tail*. The excursion is a property of the path, not a necessary cost: Section 3.1 constructs a reference path along which the past-task loss rises monotonically to its joint-optimum level (the dashed chord), so everything the accuracy curve does below its settled level and back is avoidable in principle.




Figure 1.1: 2D optimisation landscape at a task transition. Single-task losses Task A (red) and Task B (blue) are minimised at A and B ; the joint loss (purple) is minimised at $A \cap B$. The steepest-descent trajectory (green curve) illustrates the trajectory problem (Section 3.2), bowing outside A ’s basin into regions of higher loss. In contrast, the dashed black chord sketches the monotone reference path of Section 3.1, which stays inside the basin. Overshoot drivers (Section 3.2) further push the iterate off these paths. An interactive version, in which landscapes, task minima, and both gates of Section 3.3 can be manipulated, is available at https://arthurreuss.github.io/ValleyWalk/loss_landscape_visualization.html.

Level two is that real optimisers overshoot whatever path they follow, pushed by the maximally imbalanced first update at the boundary [16] and by finite-buffer gradient noise [1]. Section 3 develops both levels.

One tool: gate the new-task gradient. Sequential replay methods share a single update, a weighted combination of the current-task and replay gradients [22]. Section 3.3 generalises this: because only the ratio of the two weights matters, any reweighting of this update can be rewritten with the replay gradient restoring at full strength and a single *gate* on the new-task gradient, $d = G g_{\text{cur}} + g_{\text{rep}}$. The gate’s *domain* can be scalar (attenuate the whole push) or directional (attenuate per direction), and its *control* feedback (react to a live signal) or feedforward (commit to a schedule or precomputed filter). This thesis builds and compares the two pure designs. The first is the **λ -curriculum**, a scalar gate: instead of switching abruptly to the joint loss, it ramps the new-task loss weight $\lambda(t)$ from 0 to 1, so the minimiser of the objective moves along a continuous path from θ_A^* to θ_{joint}^* and the model tracks a moving valley floor. Its adaptive variant sets λ from the ratio of the replay and new-task gradient norms, a live damage signal, which makes it a scalar *feedback* gate. The second is **asymmetric preconditioned experience replay** (asymmetric PER), a directional *feedforward* gate: it filters the new-task gradient through the curvature of the replay loss, passing the push at full gain where the past task is flat and braking it where the past task is steep, while the restoring replay gradient stays untouched.

2 Background and Related Work

2.1 Continual Learning and Experience Replay

Methods against catastrophic forgetting fall into a few families: replay, parameter and functional regularisation, optimisation-based methods, context-dependent processing, and template-based classification [25]. This thesis focuses on replay, specifically its simplest formulation: Experience Replay (ER) [3]. At each step, ER draws a replay mini-batch from a buffer of stored past examples, combines it with the current-task mini-batch, and takes one optimiser step on the summed loss.

A standard approach to managing this episodic memory is reservoir sampling [26], which ensures that every sample seen so far occupies a buffer slot with equal probability. Under this setup, the buffer remains empty during the first task, meaning ER reduces to plain SGD; consequently, performance gaps and forgetting metrics are evaluated from the second task onward. Despite this simplicity, ER remains highly competitive: a small episodic memory often outperforms purpose-built methods and scales to longer sequences [3]. Furthermore, because the ER update consists of exactly two ingredients, it provides a remarkably clean substrate for isolating and studying interventions at task boundaries.

However, a finite buffer is inherently an imperfect estimator of the past-task distribution. Aljundi et al. [1] frame buffer selection as approximating the feasibility region of the full past data, with reservoir sampling serving as one specific policy within that framework. Consequently, the replay gradient derived from a small buffer will differ in both direction and magnitude from the true gradient of all past data.

2.2 The Stability Gap

Lange et al. [16] identified the stability gap using continual evaluation, measuring past-task accuracy after every update rather than only at task boundaries. In this view a new task begins with substantial but transient forgetting that the model subsequently recovers from, and they quantify it with worst-case metrics, mainly the minimum past-task accuracy reached during training, rather than with end-of-task averages alone. The effect is not confined to one method: they observe it across experience replay, constraint-based replay, knowledge distillation, and parameter regularisation. It also persists well beyond the standard benchmarks, appearing in joint incremental learning of homogeneous tasks [13] and in continual pre-training of large language models [7], and several mitigations have been proposed [9, 23].

Two findings in this literature anchor the account

of Section 3. First, Lange et al. [16] trace the initial drop to a gradient imbalance at the boundary: the new-task gradient dominates the first updates, so the optimiser trades stability for plasticity precisely when the past task is most exposed. Second, the gap is a property of the trajectory rather than the objective. Hess et al. [11] show it persists under training on the exact joint loss with full access to all past data, and Kao et al. [14] demonstrate the geometric mechanism analytically: steepest descent from an old-task minimum leaves the old basin even with perfect joint gradients, because the steepest direction and the direction to the joint minimum disagree. A perfect objective therefore fails to prevent the gap, and the descent geometry itself must be addressed.

2.3 One Update, Many Methods

Beyond plain ER, several methods constrain the update at the task boundary using information about past tasks. GEM stores past-task gradients and projects the new-task gradient onto the nearest direction that leaves every stored loss unchanged or lower [17]. A-GEM replaces the per-task constraints with a single averaged past-task gradient, reducing the projection to one cheap operation [2]. EWC adds a quadratic penalty around the past optimum weighted by a curvature estimate [15], and natural-gradient approaches precondition the update with past-task curvature so that steps avoid directions the old task depends on [14].

On the surface these methods look unrelated, but Sodhani et al. [22] show the replay-based ones share a single update. Every step combines two gradients, the current task’s and the replay buffer’s, under two scalar weights,

$$\theta \leftarrow \theta - \eta (\alpha_1 \nabla L_{\text{current}} + \alpha_2 \nabla L_{\text{replay}}), \quad (2.1)$$

so a method reduces to a rule for choosing α_1 and α_2 . Plain ER fixes both to one. A-GEM holds $\alpha_1 = 1$ and raises the replay weight α_2 only when the new-task and stored gradients conflict, by just enough to cancel the interference, so its projection is equivalent to re-weighting the two losses; GEM does the same per stored task, through a quadratic program whose multi-constraint solution no longer reduces to two scalar weights. MEGA [8] likewise holds $\alpha_1 = 1$ but sets the replay weight from a loss ratio, $\alpha_2 = L_{\text{replay}}/L_{\text{current}}$, leaning on memory once the new task is already fit.

The same reading applies to the λ -curriculum of this thesis. It fixes $\alpha_2 \equiv 1$ and turns the other knob, ramping the current-task weight $\alpha_1 = \lambda(t)$ up from zero. It belongs to the same family as A-GEM and MEGA; what distinguishes it is which weight it moves and what signal drives the weight. The adaptive variant (Section 3.3) is the mirror image of MEGA:

MEGA raises the *replay* weight from a *loss* ratio to recruit memory once the new task is learned, while the adaptive curriculum holds the *new task* back with a *gradient-norm* ratio, letting it in only as the replay gradient reasserts itself.

2.4 Low-Loss Paths and Mode Connectivity

Mode connectivity gives empirical backing to the *monotone reference path* introduced in Section 3: the route from the old-task solution θ_A^* to the joint solution θ_{joint}^* along which the past-task loss only ever rises toward its final value, rather than first climbing above it and coming back down as the bowed descent does. For such a path to be usable, the two solutions must be joined by a corridor of low loss in the first place, and this is what the mode-connectivity literature reports. Garipov et al. [6] and Draxler et al. [4] showed that independently trained minima of a deep network are joined by curved paths of near-constant low loss, and Frankle et al. [5] sharpened this to straight paths for networks that share early training. For continual learning specifically, Mirzadeh et al. [20] found that the sequentially trained and the multitask solutions often lie in the same basin, connected by a straight low-loss line. A low-loss route from θ_A^* to θ_{joint}^* should therefore exist in trained networks. Where it does, the stability gap reflects the optimiser wandering off that route, and an intervention succeeds by keeping it there.

3 The Stability Gap as a Trajectory Problem

This section reframes the stability gap as a two-level trajectory problem. Learning a new task incurs a permanent stability-plasticity cost; the gap is the transient excess above it. This excess has two sources: a geometric bow in the ideal descent path, and overshoot that real training dynamics add on top. We formalise the split by constructing a monotone reference path, then show that a single design choice, how the update gates the new-task gradient, controls both sources. Two contrasting gates instantiate it: an adaptive scalar curriculum and an asymmetric directional filter.

3.1 The Reference Path and the Tracking Bound

Learning a second task inevitably moves the model. In a quadratic approximation with task losses L_A and L_B and Hessians H_A and H_B , the displacement $d = \theta_{\text{joint}}^* - \theta_A^*$ from the old optimum to the joint minimum permanently raises the replay loss by roughly

$\frac{1}{2} d^\top H_A d$. Every method pays this price; it is distinct from the stability gap.

Yet this rise can be strictly monotone. Consider the curriculum objective $L_\lambda = L_{\text{replay}} + \lambda L_{\text{current}}$; at $\lambda = 1$ it is the joint loss with the two tasks weighted equally, the convention used throughout. As λ grows from 0 to 1, the minimiser traces a *reference path* $\Theta^*(\lambda)$ from θ_A^* to θ_{joint}^* . Under local positive-definiteness, an implicit-differentiation argument along the minimiser branch (Appendix D) guarantees that the replay loss climbs directly to its final level with no excursion above it:

$$\frac{d}{d\lambda} L_{\text{replay}}(\Theta^*(\lambda)) \geq 0 \quad \text{for all } \lambda \in [0, 1]. \quad (3.1)$$

This guarantee is local: it holds only while the combined Hessian $H_A + \lambda H_B$ stays positive-definite along the tracked branch of minima. The condition is automatic near θ_A^* and persists until the new task’s negative curvature overpowers the old task’s; on a deep non-convex backbone it therefore serves as motivation rather than a guarantee.

Figure 3.1 shows this: as λ grows, the loss contours deform continuously from the old-task basin to the joint basin and the minimiser slides along the connected valley floor. Where the condition holds, the replay loss climbs directly to its final level, and any transient excess *above* it, the stability gap, is a preventable deviation from the reference path.




Figure 3.1: Evolution of the curriculum loss $L_\lambda = A + \lambda B$. As λ increases, the contours morph continuously from the old-task basin around A to the joint basin around $A \cap B$, and the current minimiser (*Curr*) slides along the moving valley floor, forming the reference path $\Theta^*(\lambda)$. At the minimiser the replay (red) and current-task (blue) contributions cancel by the first-order condition $\nabla L_{\text{replay}} + \lambda \nabla L_{\text{current}} = 0$. Frames from the interactive visualisation of Figure 1.1.

A real iterate cannot sit on this path exactly, but it can stay close to it. Introducing the new task gradually over an N -step ramp ($\lambda \approx 1/N$) keeps the iterate tethered to the moving valley floor $\Theta^*(\lambda(t))$ rather than letting it leap across the landscape; a gradient-flow analysis bounds the resulting lag by $O(1/N)$ (Appendix D). Because the monotone path fixes the endpoint at the permanent level, this caps the highest transient replay loss of the whole transition,

$$\max_t L_{\text{replay}}(\theta(t)) \leq L_{\text{replay}}(\theta_{\text{joint}}^*) + O\left(\frac{1}{N}\right), \quad (3.2)$$

the first term the permanent price, the second a transient that shrinks with ramp length. A slow ramp is precisely what the λ -curriculum of this thesis implements.

3.2 The Two-Level Gap: Bowing and Overshoot

Real trajectories deviate from the reference path on two distinct levels.

Level One (The Bow): Even under idealised conditions (exact gradients, vanishing step size), the steepest-descent direction of the abruptly assembled joint loss misaligns with the direction toward the joint minimum whenever the loss curves more sharply in some directions than in others. The trajectory bows outward, leaving the old task’s basin before returning [14]. This deterministic out-and-back arc forms the structural core of the gap (Figure 1.1, green curve).

Level Two (Overshoot): Real training dynamics inflate this arc. At a sudden task boundary, the minimum teleports from θ_A^* to θ_{joint}^* , leaving the iterate exposed to a large new-task gradient while the replay gradient is near zero [16]. The acute driver is this imbalance met by a finite step size: the first discrete steps ride the unopposed push far out of the old basin, creating the sharp first-step drop. A chronic driver, estimator variance from finite buffer sampling (Section 2.1), injects noise into the restoring gradient at every step; it slows the recovery tail and can cause sudden drops long after the boundary.

The two levels shape the gap’s two dimensions: depth records the overshoot, whether the acute boundary spike or a late noise-driven drop, while the hanging tail is the Level-One arc plus the slow, noisy recovery the chronic driver leaves behind. This yields the falsifiable core, formalised as RQ1 in Section 4: removing the overshoot drivers should flatten the drop and leave the arc, placing the optimiser on the reference path should remove the arc and leave the overshoot, and doing both should close the gap while the permanent cost of Section 3.1 remains.

3.3 Two Gates: Attenuating the New-Task Gradient

Equation (2.1) shows that replay methods differ mainly in how they weight the replay and current-task gradients. Because a common scale of the two weights can be absorbed into the learning rate (exactly so in gradient flow, and up to a learning-rate retune at finite step size), only their ratio matters, so any such method can be rewritten with the replay weight normalised to one. The replay gradient is the only restoring force pulling the iterate back to the reference path, so a protective method must attenuate the new-task push, the source of the bow and overshoot, while leaving the restoring gradient intact.

Generalising the current-task weight to an operator G yields a one-sided *gate*:

$$d = G g_{\text{cur}} + g_{\text{rep}}. \quad (3.3)$$

The curriculum of Section 3.1 realises this gate along one axis, *time*: it introduces the new task slowly, so the push is never large enough at once to knock the iterate off the reference path. A complementary axis is *direction*. At any single step only part of the push is harmful, because the past task is flat in most directions and steeply curved in a few, and only motion along the steep directions raises the replay loss. Preconditioning the push by past-task curvature exploits this, passing it at full gain where the past task is flat and shrinking it where the past task is steep, so retention costs little plasticity. This natural-gradient idea underlies Natural Continual Learning [14] and, in a replay setting, Online Curvature-Aware Replay [24]. Both, however, precondition the *entire* update, so the same metric that brakes the harmful push also brakes the restoring replay gradient, protecting the past task only by slowing the return to it. Asymmetric PER keeps the gate one-sided: it filters the new-task gradient alone and leaves the restoring gradient at full strength. We include the symmetric preconditioner, which damps both, as a control (Appendix C.5).

As detailed in Table 3.1, gates vary by their *domain*, scalar or directional, and their *control*: a feedback gate reacts to a live damage signal, while a feedforward gate commits in advance to a schedule or a precomputed filter. This thesis explores the two pure corners of this design space; a fixed linear ramp, the curriculum’s feedforward variant, serves as an ablation:

Adaptive λ -curriculum (Scalar Feedback): A scheduling parameter $\lambda(t)$ sets the new-task weight from the ratio of replay to new-task gradient norms, a live damage signal that fires whichever direction harm arrives from. By scaling the loss itself, it deforms the objective into the homotopy of Section 3.1, so the optimiser’s natural descent follows the reference path.

Asymmetric PER (Directional Feedforward): The gate $G = \delta(F_{\text{rep}} + \delta I)^{-1}$ is a fixed filter set by the replay Fisher F_{rep} , giving the update

$$d = \delta(F_{\text{rep}} + \delta I)^{-1} g_{\text{cur}} + g_{\text{rep}}, \quad (3.4)$$

with the damping δ its single knob: as $\delta \rightarrow 0$ the push concentrates in the directions the past task is flat in. Committed in advance, it is feedforward.

The taxonomy yields the hypotheses that Section 4 formalises as research questions. The scalar feedback gate holds back the whole push and so pays a plasticity cost a separate accelerator must repay. The directional feedforward gate should protect the past task

Method	Gate G	Domain	Control
Vanilla ER	I	—	—
A-GEM	reweight on conflict	scalar	feedback
MEGA	loss ratio	scalar	feedback
Linear curric.	$\lambda(t) I$, clock	scalar	feedforward
Adaptive curric.	$\lambda(t) I$, norm ratio	scalar	feedback
Asym. PER	$\delta(F_{\text{rep}} + \delta I)^{-1}$	directional	feedforward

Table 3.1: Replay methods written as gates on the new-task gradient, Equation (3.3): each method’s gate G , its domain (scalar or directional), and its control (feedback or feedforward). The two methods of this thesis occupy the pure corners of the two axes: a scalar feedback gate and a directional feedforward gate.

at little plasticity cost. The corners also differ in computational cost: the scalar gate adds only a gradient-norm ratio per step, while the directional gate solves a linear system in the replay Fisher at every step, by conjugate gradients over Fisher-vector products (Appendix A.4).

Momentum is the natural accelerator, and it should discriminate between the corners by *where* each gate attenuates. The curriculum attenuates at *source*: with $\lambda \approx 0$ early there is nothing to accumulate, so momentum composes with the gate and repays its plasticity debt; the preconditioner attenuates each *step* while the gradient keeps its size, so accumulated velocity can rebuild the displacement the filter removes and should partly undo the protection. On vanilla ER, whose unopposed first step carries no accumulated velocity, momentum should leave the depth unchanged, so the head-to-head comparison of the corners should depend on the momentum regime rather than trade along a fixed stability–plasticity frontier.

4 Research Questions

The two-level account and tracking bound of Section 3 make falsifiable predictions about the gap and each gate. The experiments are organised around four questions. The first three are answered on rotated MNIST with a small MLP, where the gap can be probed step by step across many conditions and seeds; the fourth asks whether those answers survive a deeper backbone and a stronger task shift.

RQ1: The account. Does the gap separate into a deterministic arc of the descent path and overshoot of the followed path, such that removing the overshoot drivers leaves the arc, placing the optimiser on the reference path leaves the overshoot, and doing both removes the gap while the permanent stability–plasticity cost stays visible as reduced plasticity?

RQ2: The scalar feedback gate. Does the λ -curriculum reduce the gap by timing the new-task

push, and how do its ingredients divide the work? Four sub-questions probe the design:

RQ2a (linear ramp). Does gap depth fall as the ramp length N grows, consistent with the $O(1/N)$ bound?

RQ2b (adaptive schedule). Does the gradient-norm feedback signal match the best fixed ramp while removing the knob N ?

RQ2c ($\lambda_{\min}$ floor). Does a small floor $\lambda_{\min} > 0$ recover early plasticity at a modest depth cost?

RQ2d (momentum). Does momentum compose with the curriculum, closing the stochastic residual and restoring final accuracy, while leaving the depth of vanilla ER unchanged?

RQ3: The directional feedforward gate. Does asymmetric PER reduce the gap as its filter strengthens, at what plasticity cost, and where does a feedforward gate fail? The account predicts one failure axis, momentum, whose accumulated velocity rebuilds the push the filter shrinks one step at a time.

RQ4: Generalisation. Do the findings carry to a deeper backbone (ResNet-18) and a stronger task shift (CIFAR-10 domain corruption): does each gate still cut the gap at matched final accuracy, and do the momentum signatures of RQ2d and RQ3 reproduce there in a momentum cross?

5 Methods

This section describes the empirical testbeds (Section 5.1), the metrics that summarise each run (Section 5.1.2), and the experimental blocks (Section 5.2) that answer the research questions of Section 4. All code, experiment configurations, and analysis scripts are available at <https://github.com/Arthurreuss/ValleyWalk>.

5.1 Experimental Setup

5.1.1 Testbeds

Two testbeds of different scale carry the study. A small, fast one (rotated MNIST on a two-layer MLP) runs the dense mechanistic sweeps: the driver block, the curriculum sweep, the asymmetric-PER damping sweep, and the momentum cross. A larger one (domain-shift CIFAR-10 on a CIFAR-adapted ResNet-18) tests whether the findings survive a deeper backbone and a stronger task shift, the regime where the local condition behind the bound of Section 3.1 weakens. Both hold the label space and task structure fixed and shift only the input distribution, and both use two tasks in their main blocks, so there is exactly one transition and one instance of the stability

gap; a five-task rot-MNIST sequence and a three-task CIFAR sequence serve as multi-boundary checks.

On rot-MNIST, following Lange et al. [16], T_0 presents the digits at -90° and T_1 applies a further 90° rotation, each task using the standard 60,000 training and 10,000 test images. On CIFAR-10, T_0 is clean and T_1 applies Gaussian-noise corruption at severity 3 [10]; the ten classes are shared, so only the input distribution shifts. A three-task variant ([none, Gaussian, shot]) and two alternative corruptions appear in the generalisation block. The rot-MNIST MLP is cheap enough to sweep across many conditions and seeds and is where the local condition of Section 3.1 is most plausible; the ResNet-18 is the deeper, more realistic counterpart. rot-MNIST runs online (one epoch per task); the ResNet needs several passes to converge before the transition, so the CIFAR block uses ten epochs per task. Optimiser, learning rate, and batch size are common across benchmarks. Full backbone and protocol specifications are in Appendix A.2 (Tables A.1–A.3).

Throughout training every task’s test accuracy is evaluated every 10 steps; at each transition the cadence rises to every step, so the dip is recorded at full resolution (the whole online epoch on rot-MNIST, a 250-step window on CIFAR). A stability-gap tracker snapshots each past task’s pre-switch accuracy and records the resulting (step, accuracy) series. For the driver block only, a gradient tracker additionally records the fidelity of the replay-buffer gradient against the exact past-task gradient, at the same cadence; each such step costs a full pass over the past-task training set, which is why it is not run elsewhere.

5.1.2 Metrics

Each metric is reported as the seed mean $\pm$ sample standard deviation; complete definitions are in Appendix B. The gap definition of Section 3, the transient excess below the level a task settles at, gives two primary metrics. With a_j^{pre} the pre-switch baseline on past task j and $a_j(t)$ its accuracy at step t , **gap_depth** is the maximum instantaneous excursion $\max_j(a_j^{\text{pre}} - \min_t a_j(t))$, which Section 3.2 predicts to be overshoot-dominated. The second, **gap_area_end**, integrates the shortfall below the settled level: with reference $b_j = \min(a_j^{\text{pre}}, \bar{a}_j)$, where $\bar{a}_j$ is the accuracy the task eventually reaches, it sums $\max(0, b_j - a_j(t))$ over the new task’s training steps and across past tasks. Anchoring to the recovered level implements the split of Section 3.1: a permanent step down to a lower plateau contributes to the standard forgetting metrics and near-zero to the area, while a true dip-and-recover registers in full. One consequence, revisited when reading the tables, is that two methods settling at different levels are measured against different lines, so a better-recovering method can carry the *larger* area; a positive area marks transient disruption,

not net cost, and we read it with final accuracy. The arc of Section 3.2 and the slow part of the noisy recovery live here; sharp noise-driven drops register in depth instead. A third metric, **recovery_steps**, and the gradient-fidelity diagnostics are defined in Appendix B.

To place the gap in the standard frame we report the continual-evaluation metrics of Lange et al. [16] from the task-boundary accuracy matrix: average final accuracy (ACC), which checks that a gate lowers the gap at preserved accuracy; worst-case retention (min-ACC, WC-ACC); and short-horizon plasticity (WP_{100}). FORG and the windowed WF_w/WP_w family are logged for completeness. Conditions are compared with the paired Wilcoxon signed-rank test on seed-paired differences, preferred over the t -test because five paired seeds cannot support a normality assumption on bounded metrics: the gap metrics floor at zero, where the strongest conditions cluster, and accuracy sits near its ceiling, so the seed differences are skewed rather than Gaussian. With five paired seeds the smallest two-sided p -value is 0.0625, reached when all five differences share a sign; we read $p = 0.0625$ as complete directional agreement and let effect sizes separate large results from marginal ones.

5.2 Experiments

Each block maps to one or more research questions and tests the corresponding predictions of Section 3. All conditions use five seeds and the protocol above unless noted.

5.2.1 Overshoot Drivers (RQ1)

The driver block answers the overshoot side of RQ1: each condition switches off one of the drivers of Section 3.2 by construction and observes what remains. The complementary path-side intervention is the curriculum, whose full-buffer diagnostic below completes the RQ1 cross. Table 5.1 lists the four conditions. Balancing (D3, D4) gives g_{new} and g_{replay} equal magnitude from the first step, switching off the boundary imbalance: an unbalanced step at θ_0^* (the converged T_0 optimum, θ_A^* of Section 3) is driven almost entirely by g_{new} , a balanced step carries the replay direction at equal weight. The full-data buffer (D2, D4) replaces vanilla ER’s noisy 256-sample replay gradient with the exact empirical past-task gradient, switching off the estimator variance (mechanics in Appendix A.3). Because the drivers interact (Section 3.2), we read the conditions as interventions, not an additive decomposition: depth differences measure how much of the drop each driver carries in that context, while the area and per-step shape carry the level-one signal: with the replay estimate exact (D2, D4), a dip-and-recover bend that remains cannot come from the variance driver, the arc’s fingerprint.

Label	Condition	Replay buffer	Gradient handling		Driver removed by construction
D1	Vanilla ER	1k reservoir	unbalanced		none (all drivers active)
D2	Full-data ER	60k (all past)	unbalanced		estimator variance (exact replay gradient)
D3	Balanced ER	1k reservoir	unit-norm	balanced	boundary imbalance
D4	Full-data + balanced	60k	unit-norm	balanced	imbalance and variance

Table 5.1: Driver-removal conditions. Each switches off one or two overshoot drivers relative to vanilla ER; what all four leave in place is the descent path itself, so the arc should survive throughout.

5.2.2 The λ -Curriculum Sweep (RQ2, RQ1)

The curriculum sweep tests the scalar feedback gate one ingredient at a time: the $O(1/N)$ prediction of Equation (3.2) (RQ2a), the adaptive schedule (RQ2b), and the $\lambda_{\min}$ floor (RQ2c); a full-buffer diagnostic (C8) then returns to RQ1, testing the claim that the curriculum removes the arc and leaves only stochastic overshoot. All conditions run on top of standard ER (1k reservoir, unbalanced gradients), so the schedule is the only changed variable (Table 5.2). The adaptive schedule (C4) is the feedback variant: it sets λ to a clipped exponential moving average of the gradient-norm ratio $r := \|g_{\text{replay}}\|/\|g_{\text{new}}\|$. At θ_0^* the replay gradient is near zero by stationarity, so r and λ start near 0; as the iterate moves, $\|g_{\text{replay}}\|$ grows and $\|g_{\text{new}}\|$ shrinks, so $\lambda \rightarrow 1$ at the pace the damage signal dictates, with no ramp length to choose. The floor $\lambda_{\min}$ (C5–C7) keeps the schedule off $\lambda \approx 0$ during the first steps and buys back early plasticity.

C8 combines the practical curriculum configuration (C7) with the full 60k buffer of D2/D4, so the path is fixed and the variance driver is switched off at source. This makes it a double probe. First, it tests what C7’s remaining momentum-0 transient is made of: the curriculum has already removed the arc, so if the residual is stochastic overshoot from buffer noise, the exact replay gradient should remove it too and C8 should show no gap even without momentum. Second, comparing C8 with and without momentum (C8 vs. C8_M) separates the two roles momentum could play for the curriculum. The full buffer already removes the gradient noise, so C8 has no variance left for momentum to reduce; whatever momentum still improves on top of C8 must come from its other role, accelerating new-task learning. C8 stores the whole past task and so breaks the limited-memory premise; it is a diagnostic only, and C7 remains the practical configuration.

5.2.3 Momentum Cross (RQ2d, RQ3)

Every rot-MNIST condition of the driver D-series and the curriculum C-series, and the asymmetric-PER sweep legs P1–P2 at the settings noted in Table 5.3, is run at momentum 0.0 and 0.9; momentum-on runs carry an M subscript (e.g. D1_M). The cross tests the composability prediction of Section 3.3: whether mo-

mentum helps or hurts a method’s gap should depend on where the method attenuates the push. A five-task rot-MNIST sequence (rotations $[0, 30, 60, 90, 120]^\circ$, four transitions, five seeds) repeats the three-way comparison of vanilla ER, the practical curriculum, and asymmetric PER at both momentum settings, as a multi-boundary check (Appendix C.9).

5.2.4 Asymmetric-PER δ Sweep (RQ3)

The damping sweep tests the directional feedforward gate: does the gap shrink as the filter strengthens, at what plasticity cost, and where does the gate fail? Asymmetric PER implements Equation (3.4), filtering the new-task push through the replay Fisher F_{rep} [12]. Normalising by δ keeps flat directions at unit gain: along an eigendirection of curvature ρ the gain is $\delta/(\rho + \delta)$, so directions the past task is flat in pass through whole and steep ones are shrunk toward zero. The damping δ is the single knob: at $\delta \rightarrow \infty$ plain ER returns, while as $\delta \rightarrow 0$ the push concentrates in replay-flat directions with the restoring gradient g_{rep} at full strength throughout. The push is applied matrix-free by conjugate gradients over Fisher-vector products (Appendix A.4).

Table 5.3 lays out the sweep: $\delta \in \{1.0, 0.3, 0.1, 0.03\}$ crossed with two buffer legs, each at both momentum settings. The *standard* leg (1k reservoir) matches the curriculum’s practical regime, while the *full-buffer* leg makes the replay gradient exact, switching off the variance driver as in D2/D4, so the remaining area measures the deterministic arc and the sweep reads out directly whether the gate shrinks the arc itself. Two controls guard the reading: a solver control reruns the strongest filter at 60 CG iterations to separate gate properties from an under-converged solve (Appendix A.4), and a symmetric control preconditions the summed gradient by the same replay Fisher, isolating the effect of leaving the restoring gradient untouched (Appendix C.5).

5.2.5 Generalisation: CIFAR-10 (RQ4)

The CIFAR-10 block tests whether both gates and the momentum signatures carry to a deeper backbone and a stronger task shift. All CIFAR conditions use the ResNet-18 protocol of Table A.3 (ten epochs per

Label	Condition	$\lambda(t)$ schedule	Hyperparameter	Tests
—	Vanilla ER (baseline)	$\lambda \equiv 1$ from step 1 (abrupt switch)	—	reference
C1	Linear $N=50$	$\text{clip}(t/50, 0, 1)$	$N = 50$	RQ2a
C2	Linear $N=100$	$\text{clip}(t/100, 0, 1)$	$N = 100$	RQ2a
C3	Linear $N=200$	$\text{clip}(t/200, 0, 1)$	$N = 200$	RQ2a
C4	Adaptive	$\text{clip}(\text{EMA}_\alpha(r), 0, 1)$	$\alpha = 0.05$	RQ2b
C5	Adaptive + $\lambda_{\min}=0.05$	$\max(\lambda_{\min}, \text{clip}(\text{EMA}_\alpha(r), 0, 1))$	$\alpha = 0.05, \lambda_{\min}=0.05$	RQ2c
C6	Adaptive + $\lambda_{\min}=0.10$	as C5	$\lambda_{\min}=0.10$	RQ2c
C7	Adaptive + $\lambda_{\min}=0.20$	as C5	$\lambda_{\min}=0.20$	RQ2c
C8	C7 + full 60k buffer	as C5, exact replay gradient	$\lambda_{\min}=0.20$	RQ1, RQ2d

Table 5.2: λ -curriculum conditions, all applied on top of standard ER. EMA_α is an exponential moving average with smoothing α , and $r := \|g_{\text{replay}}\|/\|g_{\text{new}}\|$ is the gradient-norm ratio. Each condition is run at momentum 0.0 and 0.9.

Label	Condition	Replay buffer	Damping δ	Purpose
P1	Standard leg	1k reservoir	1.0, 0.3, 0.1, 0.03	practical regime, matched to the curriculum (RQ3)
P2	Full-buffer leg	60k (exact g_{rep})	1.0, 0.3, 0.1, 0.03	variance driver off; residual area is the arc (RQ3, RQ1)
P3	Solver control	60k	0.03	60 CG iterations instead of 25; rules out solve artefacts
P4	Symmetric control	1k reservoir	1.0, 0.3, 0.1, 0.03	preconditions $g_{\text{cur}} + g_{\text{rep}}$; ablates the one-sided gate

Table 5.3: Asymmetric-PER conditions. Both sweep legs run at momentum 0.0; momentum-0.9 counterparts cover $\delta \leq 0.3$ on the standard leg and $\delta \in \{0.3, 0.1\}$ on the full-buffer leg. The two controls run at momentum 0, where the effects they check for arise. The filter metric is always the replay Fisher F_{rep} , estimated on a sampled batch (Appendix A.4).

task, buffer 5,000). The headline conditions run at momentum 0.9 (Table 5.4, top), each method crossed with BatchNorm and GroupNorm because the two respond differently at the switch: the input shift disrupts BatchNorm’s running statistics, which must re-adapt during the first steps, while GroupNorm uses no cross-batch statistics, so its transient reflects optimisation dynamics alone. The curriculum runs as the adaptive schedule at $\lambda_{\min} = 0.20$; asymmetric PER runs at $\delta = 0.1$, a strong rot-MNIST setting. Both carry their rot-MNIST configurations without a CIFAR-specific sweep, so they meet the deeper backbone on equal footing. A momentum cross (Table 5.4, bottom) adds momentum-0 runs of all three methods at GroupNorm, repeating the three-way composability test where the transient is free of normalisation-statistics effects; because momentum also drives convergence on this backbone, accuracy is compared within a momentum setting. A further generalisation block pairs the practical curriculum against a paired vanilla-ER control on shifts it was never tuned on (a two-task shot-noise shift, a two-task contrast shift, and a three-task [none, Gaussian, shot] sequence, each at both normalisations over three seeds) to test whether the result depends on the particular corruption or on having a single transition.

Label	Method	Norm	μ
G1	Vanilla ER	BatchNorm	0.9
G2	Adaptive $\lambda_{\min}=0.20$	BatchNorm	0.9
G3	Asym. PER $\delta=0.1$	BatchNorm	0.9
G4	Vanilla ER	GroupNorm	0.9
G5	Adaptive $\lambda_{\min}=0.20$	GroupNorm	0.9
G6	Asym. PER $\delta=0.1$	GroupNorm	0.9
G7	Vanilla ER	GroupNorm	0.0
G8	Adaptive $\lambda_{\min}=0.20$	GroupNorm	0.0
G9	Asym. PER $\delta=0.1$	GroupNorm	0.0

Table 5.4: CIFAR-10 conditions (two-task Gaussian-noise shift, buffer 5,000, five seeds). Top: the headline block at momentum 0.9. Bottom: the momentum cross, pairing momentum-0 runs of all three methods against their momentum-0.9 counterparts at GroupNorm. The generalisation block (novel corruptions and the three-task sequence, three seeds) is described in the text and reported in Appendix C.7.

6 Results

The results follow the research questions: the two-level account (Section 6.1, RQ1), the scalar feedback gate (Section 6.2, RQ2), the directional feedforward gate (Section 6.3, RQ3), a synthesis of the two momentum crosses (Section 6.4), and the CIFAR-10 generalisation of both gates (Section 6.5, RQ4).

Reading the tables. Depth is in percentage points (pp): a depth of 19.5 means the worst past-task evaluation falls 19.5 points below its pre-switch baseline. Section 3.2 predicts depth to be overshoot-dominated, and both overshoot drivers express in it: the acute boundary spike, and the buffer’s noise, whose sudden drops can set the worst evaluation long after the switch. The end-referenced area `gap_area_end` (accuracy-steps) integrates the excursion below the level each task settles at; its reference convention, and the caveat that cross-condition areas must be read together with final accuracy, are stated once in Section 5.1.2 and apply throughout. Values are mean $\pm$ standard deviation over five seeds (three for the CIFAR generalisation block). Statistical comparisons follow the Wilcoxon convention of Section 5.1.2: a p -value is quoted only where seeds disagree, and otherwise we report complete directional agreement. The full metric set for every condition is in Appendix C.

6.1 The Two Levels, Separated (RQ1)

RQ1 asks whether the gap separates into a deterministic arc of the descent path and overshoot of whatever path is followed, such that removing the overshoot drivers and fixing the path are separately insufficient but jointly sufficient. Crossing the two interventions at momentum 0 gives the 2×2 of Figure 6.1: colour encodes the path (red: bowed steepest descent, no gate; blue: the curriculum’s valley-floor path, which tracks the reference path of Section 3.1), linestyle the overshoot drivers (solid: present; dashed: removed), so each curve is one cell of the cross.

Removing the drivers leaves the arc. Following the driver ladder (D1–D4, full metrics in Appendix C.1, Table C.1), unit-norm balancing (D3) removes the boundary imbalance and takes 13.8 pp off the depth ($19.5 \rightarrow 5.7$), by far the largest single drop; swapping the reservoir for the exact past-task gradient (D4) removes the estimator variance and takes off a further 3.8 pp ($5.7 \rightarrow 1.9$). Gradient-fidelity diagnostics confirm the driver each condition switches off: the finite-buffer D1 and D3 reach only $\text{c}\ddot{\text{o}}\text{s}(g_{\text{replay}}, g_{\text{true}}) \approx 0.71$ and 0.66 with per-step minima near zero, while the full-data D2 and D4 sit at 1.000 by construction (Appendix C.3). One nuance keeps the ladder honest: balancing is not a pure imbalance intervention, and not a gate in the sense of Section 3.3, because it rescales the replay gradient as well (unit-normalising the near-zero restoring force at the boundary amplifies it), while also stripping the magnitude component of the buffer noise and capping how far any single bad replay estimate can move the iterate. Balancing does not remove the variance driver, however: D3’s replay gradient is still the noisy 1k estimate, and its direction is as unreliable as D1’s (the mean cosine is in fact slightly lower, 0.66 vs. 0.71);

only its magnitude is controlled. Because the interventions overlap in this way, the ladder is read as a set of interventions rather than an additive decomposition of the gap (Section 5.2.1). What no driver removal touches is the descent path, and the area shows it: with the spike gone at D4 (only 1.9 pp) a clear bend remains, `gap_area_end` = 3.2. Because the full-data conditions carry no noise in the restoring gradient, that area cannot be a residual of the variance driver; it is the deterministic arc of Section 3.2. Figure 6.1 shows it directly: with the drivers removed, D4 (dashed red) traces a smooth dip-and-recover with no jitter, the fingerprint the account attaches to a bowed path rather than a spike. Because balancing modifies the vector field, D4’s bend is strictly the arc of the balanced dynamics; the unbalanced full-data D2, whose restoring gradient is exact but whose new-task push is still a stochastic mini-batch gradient, shows the same dip-and-recover bend beneath its boundary spike and residual jitter, so the arc is not an artefact of the balancing. (That D2’s area, 8.1, exceeds D1’s 6.3 is the end-referencing convention at work: D2 recovers to a higher level, so more of its excursion counts as transient; Section 5.1.2.) The full ladder, including the intermediate D2 and D3, is plotted in Appendix C.1 (Figure C.1).

Fixing the path leaves the overshoot. The other row holds the drivers in place and fixes the path with the curriculum. At momentum 0 the practical curriculum (C7, solid blue) still shows 13.9 pp of depth, and Figure 6.1 shows where it comes from: the boundary spike is gone, and the depth arrives instead as sudden noise-driven drops well after the switch, the chronic driver expressing in depth as Section 3.2 predicts. With the arc removed, essentially all of C7’s gap is overshoot. That this residual is estimator variance rather than anything structural is settled by C8, which adds the exact full buffer to C7 so that the path is fixed *and* the variance driver is switched off at source. The gap then collapses to 0.2 ± 0.3 pp of depth and 0.05 of area, the smallest depth of any condition in the study, and the C7→C8 area drop ($2.7 \rightarrow 0.05$) holds across all five seeds. Overshoot is real and untouched by the path fix; the arc is real and untouched by the driver removal; only doing both closes the gap.

The permanent cost stays on the ledger. Closing the transient does not make the transition free. C8’s final accuracy is 0.818, the lowest of the four cells and 0.055 below vanilla ER, because holding the new-task push down at momentum 0 slows the new task without an accelerator to repay it. This is the permanent stability–plasticity cost the account predicts to remain once the transient is removed; Section 6.2 shows momentum paying it back. A second, independent line of evidence that the arc is deterministic geometry comes from the directional gate (Section 6.3):




Figure 6.1: The RQ1 2×2 as per-step curves: past-task (T_0) accuracy across the rot-MNIST switch at momentum 0; seed mean with $\pm$ std bands, $x=0$ the task switch. Colour crosses the path (red: bowed steepest descent, no gate; blue: the curriculum’s valley-floor path), linestyle the overshoot drivers (solid: present; dashed: removed by balancing and the exact full-buffer replay gradient). Cell metrics as gap_depth (pp) / gap_area_end / ACC: D1 19.5/6.3/0.873, D4 1.9/3.2/0.869, C7 13.9/2.7/0.837, C8 0.2/0.05/0.818. Removing the drivers alone (D4) leaves the smooth deterministic arc; fixing the path alone (C7) leaves the overshoot, arriving as sudden noise-driven drops well after the switch; doing both (C8) holds T_0 flat through the switch, while the permanent stability–plasticity cost stays visible as the lowest final accuracy of the four.

on its full-buffer legs, where the variance driver is off and the area *is* the arc, the area shrinks monotonically as the filter strengthens ($2.3 \rightarrow 1.0 \rightarrow 0.4$ for $\delta = 1.0 \rightarrow 0.3 \rightarrow 0.1$). The arc responds to path-level intervention and to nothing else, whether the intervention times the push (the curriculum) or filters its direction (the gate).

6.2 The Scalar Feedback Gate (RQ2)

Table 6.1 reports the curriculum sweep at both momentum settings against vanilla ER. The momentum-off leg tests the schedule in isolation; the momentum-on leg is the practically relevant regime.

Lengthening the ramp reduces gap depth (RQ2a–c). At momentum 0, mean gap depth falls monotonically as the linear ramp lengthens, as the $O(1/N)$ bound of Equation (3.2) predicts (RQ2a), though successive lengths lie within the seed spread and the reduction is bought at a real accuracy cost: by $N = 200$ ACC has fallen from vanilla’s 0.873 to 0.828 and the tail has stretched, **recovery_steps** growing from 36 at $N=50$ to 133 at $N=200$, against vanilla’s ≈ 4 (Appendix C). The depth that remains, still in the mid-teens, is the untouched chronic driver: the schedule times the push but leaves the finite-buffer

replay estimate exactly as noisy as vanilla’s, so the worst evaluation keeps arriving as a late noise-driven drop, a floor no ramp length can pass and only the exact-gradient diagnostic removes (C8, Section 6.1). The adaptive schedule (C4) reaches 12.8 pp against the best linear ramp’s 14.3 without a ramp length to set (RQ2b), though the seeds are split on this difference ($p = 0.44$): at momentum 0 the adaptive signal removes the knob without losing ground, rather than demonstrably beating the best ramp. Raising the floor $\lambda_{\min}$ to 0.20 (C7) shifts the schedule in the predicted Pareto direction only weakly at momentum 0, because early T_1 progress is then set by the large deterministic boundary gradient the small floor cannot redirect (RQ2c). The momentum-0 curves of Figure 6.2 (dashed) show the trade directly: the curriculum reaches a shallower past-task dip than vanilla (a) but pays for it in slower early T_1 learning (b).

Momentum pays the plasticity debt, and closes the gap (RQ2d).

The momentum-on leg is the scalar gate’s intended regime, and momentum acts through two roles at once. Adding it to the adaptive schedule drops depth to 2.5 pp ($C4 \rightarrow C4_M$), and to 3.7 pp at the $\lambda_{\min} = 0.20$ floor ($C7_M$), and collapses the area to near zero; every depth and area contrast holds across all five seeds (Table 6.1). Accuracy, how-

Condition	Momentum 0.0			Momentum 0.9		
	gap_depth (pp)	gap_area_end	ACC	gap_depth (pp)	gap_area_end	ACC
Vanilla ER (baseline)	19.5 $\pm$ 4.1	6.3 $\pm$ 0.5	0.873 $\pm$ 0.005	15.9 $\pm$ 2.6	1.0 $\pm$ 0.2	0.928 $\pm$ 0.003
C1: Linear $N=50$	17.5 $\pm$ 1.7	4.9 $\pm$ 0.5	0.851 $\pm$ 0.023	6.4 $\pm$ 0.5	0.2 $\pm$ 0.1	0.932 $\pm$ 0.003
C2: Linear $N=100$	16.2 $\pm$ 4.3	3.6 $\pm$ 0.3	0.844 $\pm$ 0.025	6.4 $\pm$ 0.6	0.2 $\pm$ 0.1	0.932 $\pm$ 0.003
C3: Linear $N=200$	14.3 $\pm$ 4.6	1.8 $\pm$ 0.6	0.828 $\pm$ 0.025	6.2 $\pm$ 0.6	0.1 $\pm$ 0.0	0.925 $\pm$ 0.007
C4: Adaptive	12.8 $\pm$ 2.9	2.2 $\pm$ 0.5	0.833 $\pm$ 0.025	2.5 $\pm$ 0.5	0.3 $\pm$ 0.1	0.917 $\pm$ 0.001
C5: Adaptive $\lambda_{\min}=0.05$	13.2 $\pm$ 3.2	2.2 $\pm$ 0.5	0.835 $\pm$ 0.025	2.6 $\pm$ 0.5	0.1 $\pm$ 0.0	0.922 $\pm$ 0.001
C6: Adaptive $\lambda_{\min}=0.10$	13.9 $\pm$ 3.3	2.4 $\pm$ 0.5	0.836 $\pm$ 0.026	3.0 $\pm$ 0.4	0.0 $\pm$ 0.0	0.927 $\pm$ 0.002
C7: Adaptive $\lambda_{\min}=0.20$	13.9 $\pm$ 3.4	2.7 $\pm$ 0.6	0.837 $\pm$ 0.031	3.7 $\pm$ 0.4	0.0 $\pm$ 0.0	0.931 $\pm$ 0.002

Table 6.1: λ -curriculum sweep on rot-MNIST against vanilla ER, at momentum 0.0 and 0.9; five seeds. The scalar feedback gate reaches a few points of depth at matched accuracy only in the momentum-on regime it is intended for.




Figure 6.2: The scalar feedback gate and momentum’s two roles: per-step accuracy across the rot-MNIST switch, seed mean with $\pm$ std bands, for vanilla ER (red) and the practical adaptive curriculum $\lambda_{\min}=0.20$ (blue), each at momentum 0.9 (solid) and 0.0 (dashed). (a) Stability (T_0): only curriculum-with-momentum (solid blue) holds the past task near its plateau through the switch; momentum barely dents vanilla’s spike, and both methods open a deep, jittery gap without it (dashed). (b) Plasticity (T_1): at momentum 0 the curriculum slows early new-task learning (dashed blue), the plasticity debt that momentum pays back (solid blue).

ever, returns to vanilla ER’s level only together with the floor: $C4_M$ stays 1.1 pp below vanilla (0.917 vs. 0.928) while $C7_M$ matches it (0.931), so the adaptive signal alone answers RQ2b on depth, and recovering vanilla accuracy takes RQ2c’s floor as well. The depth and area collapse is variance reduction: momentum averages the noisy finite-buffer replay gradient across accumulated velocity, damping the chronic driver once the schedule has tempered the acute boundary push at source. The accuracy restoration is the second, separate role: acceleration, driving the new task hard enough to repay the plasticity the schedule borrows.

The floor, weak at momentum 0, becomes a clean Pareto knob once momentum is active, and what it trades is the permanent stability–plasticity split rather than transient stability. In this regime the past task no longer dips and recovers; it settles at a marginally lower plateau (area ≈ 0), so depth here reads where between the two tasks the shared pa-

rameters come to rest. Raising the floor from 0 to 0.20 ($C4_M \rightarrow C7_M$) shifts that split toward the new task. The final per-task accuracies (from the logged accuracy matrix, Appendix B) show both ends of the trade: T_0 settles at 0.921 instead of 0.936, while T_1 finishes at 0.941 instead of 0.899, a gain of 4.2 pp on the new task for 1.5 pp of the old, both trends across all five seeds. The gain is already visible at short horizon (WP₁₀₀; Appendix C), because a higher floor lets T_1 train from the first post-switch step. $C7_M$ is the configuration carried to CIFAR.

That momentum here closes a gap it left untouched under vanilla ER is the composability that Section 3.3 predicts for a gate that attenuates at source, and a full-buffer diagnostic separates momentum’s two roles cleanly (Table C.3, Appendix C.2): with the full buffer standing in for momentum’s averaging, C8 already has essentially no gap at momentum 0 (0.2 pp), so the residual C7 leaves is variance, closable by source

variance reduction rather than by acceleration; adding momentum to C8 then lifts ACC to 0.955 while re-introducing only a 2.0 pp transient, a brief boundary dip traced to momentum amplifying the $\lambda_{\min}$ -floored share of the push, the one share the source gate does not attenuate (Appendix C.2). The headline is that *the curriculum needs momentum for accuracy, while its residual gap needs only variance reduction, which the full buffer supplies as well as momentum does*. The full 60k buffer defeats the limited-memory premise and is diagnostic only; the practical configuration is C7_M (3.7 pp / 0.03 / 0.931) at 1/60th the memory.

For completeness the cross also confirms the vanilla-ER null: momentum shifts vanilla depth only from 19.5 to 15.9 pp, smaller than the seed spread with seeds split on its sign ($p = 0.44$), because the deterministic unopposed first step carries no accumulated velocity to average; it does cut vanilla’s tail (area $6.3 \rightarrow 1.0$, all seeds). Momentum smooths vanilla’s recovery, but it leaves the depth, the gap itself, intact.

6.3 The Directional Feedforward Gate (RQ3)

RQ3 asks whether asymmetric PER shrinks the gap as its filter strengthens ($\delta \rightarrow 0$), at what plasticity cost, and where a feedforward gate fails. Table 6.2 reports the δ sweep on both buffer legs at momentum 0; the full metric set at both momentum settings is in Appendix C.4. The standard 1k leg is the gate’s practical configuration; the full-buffer leg plays the same diagnostic role as C8, making the replay gradient exact so that the remaining area is the arc, and it does not enter the head-to-head comparisons.

Monotone, and at no plasticity cost. On the standard buffer, depth falls from 9.0 to 4.8 pp as δ goes from 1.0 to 0.03 while ACC *rises* from 0.869 to 0.880 and WP₁₀₀ from 0.733 to 0.741; every one of these trends holds across all five seeds (Figure 6.3(a)). The roughly 4 pp of accuracy the scalar gate pays at momentum 0 never materialises here, because the filter blocks the push only in the directions the past task is steep in and lets it through everywhere else. This is the scalar gate’s weakness inverted, and it is what makes the two gates complements. The full-buffer leg isolates the arc: with the variance driver off, the area shrinks with the filter, $2.3 \rightarrow 1.0 \rightarrow 0.4$ for $\delta = 1.0 \rightarrow 0.3 \rightarrow 0.1$ (the $1.0 \rightarrow 0.3$ drop holds across all seeds), the independent level-1 evidence promised in Section 6.1: the directional gate bends the followed path back toward the reference, as the curriculum does by another route. The contrast with driver removal makes the mechanism point: D4 switches off both overshoot drivers and still carries the arc (area 3.2), while the gate on the same exact-gradient leg takes the area to 0.4 at $\delta=0.1$. The gate reaches the

arc itself, which no driver removal does. This is a mechanism reading rather than an operating point, since the full buffer that isolates the arc defeats the limited-memory premise. Leaving the restoring gradient g_{rep} at full strength is what makes this a net removal of damage rather than a reshuffling of it: a symmetric variant that preconditions the summed update strangles the restoring gradient harder than the harmful push, so its forgetting *climbs* as the filter strengthens (Appendix C.5, Table C.6).

At momentum 0 the directional gate dominates the scalar gate. Head to head at momentum 0, the strongest standard-buffer setting ($\delta=0.03$: 4.8 pp / 0.880) beats the practical curriculum (C7: 13.9 pp / 0.837) on both axes at once, lower in depth and higher in accuracy. Part of this is that the curvature filter also limits how far any single noisy gradient can move the past-task loss, so it suppresses the per-step *expression* of buffer variance, which the scalar gate can only counter with momentum or a larger buffer. The scalar gate’s case is not this corner; it is momentum composability, which the next results show the directional gate lacking, and robustness to the failure mode reported next.

The cliffs: an unexpected late failure. The one place the sweep breaks monotonicity is the strongest full-buffer setting, $\delta = 0.03$, where mean depth jumps to 9.5 pp with a std of 8.2: in 3 of 5 seeds the past-task curve holds flat for roughly two hundred steps and then slides off a late cliff, the red per-seed traces in Figure 6.3(b). A solver control rules out the most straightforward explanation: rerunning the setting with the conjugate-gradient solver at 60 iterations instead of 25 reproduces the cliffs at the same seeds and steps (11.5 pp mean depth), so they are not an artefact of an under-converged solve. The cliffs came as a surprise: the account made no prediction of a late, abrupt failure, and nothing in the rest of the sweep hints at one. They are also confined to the exact-gradient diagnostic: at the same δ on the standard 1k buffer no seed shows one (depth 4.8 ± 1.9 pp, the monotone end of the sweep). The failure is tied to the exact replay gradient: the sampling jitter of a practical buffer appears to protect against it, the same noise that elsewhere drives the overshoot here keeping the iterate from settling into the escape direction. Why the cliffs happen we can only partly say. A reading consistent with the data is a blind spot of the gate’s model of the past task: the filter passes new-task pressure through directions its sampled Fisher rates as flat, and in these seeds that trust appears to be misplaced further out along the trajectory. We offer this as a hypothesis, not an established mechanism. What the control does establish is that the strongest filter can fail late and abruptly when nothing perturbs it. Because the failure is confined to the

δ	Standard (1k) buffer				Full (60k) buffer			
	depth (pp)	area	ACC	WP ₁₀₀	depth (pp)	area	ACC	min-ACC
1.0	9.0 ± 0.6	3.3	0.869	0.733	4.0 ± 1.0	2.3	0.898	0.841
0.3	7.4 ± 1.1	2.3	0.873	0.736	1.5 ± 1.1	1.0	0.900	0.867
0.1	5.9 ± 1.5	1.6	0.876	0.739	2.2 ± 1.6	0.4	0.901	0.860
0.03	4.8 ± 1.9	1.0	0.880	0.741	9.5 ± 8.2	0.4	0.903	0.787

Table 6.2: Asymmetric-PER δ sweep on rot-MNIST at momentum 0, five seeds. Depth and area fall as $\delta \rightarrow 0$ while ACC and WP₁₀₀ rise: the directional gate protects the past task at no plasticity cost. On the full buffer, a diagnostic leg where the area is the deterministic arc, the arc shrinks with the filter, its area falling to 0.4 at $\delta=0.1$, well below what driver removal leaves in place (D4: 3.2). The far cell is the exception: at $\delta=0.03$ full buffer the mean depth jumps to 9.5 pp with a large std, the cliffs of the text.




Figure 6.3: The directional feedforward gate: per-step past-task (T_0) accuracy across the rot-MNIST switch at momentum 0, sweeping the damping δ (weaker filter light $\rightarrow$ stronger filter dark). (a) Standard 1k buffer: the boundary dip shrinks monotonically as $\delta \rightarrow 0$, and accuracy rises with it (Table 6.2), so the protection is bought at no plasticity cost. (b) Full 60k buffer (variance driver off, so the residual excursion is the deterministic arc): the arc shrinks with the filter for $\delta \geq 0.1$, but at the strongest setting $\delta=0.03$ the gate fails late and abruptly: shown per seed (critical red), 3 of 5 seeds hold flat for ~ 200 steps and then slide off a late cliff, a failure mode a feedback signal does not share.

diagnostic leg, it is unlikely in a practical configuration, but it exposes a real weakness of feedforward control.

Momentum re-inflates the per-step filter. Momentum undoes the directional gate’s protection, the opposite of what it does for the scalar gate, and it does so on both legs: on the practical 1k buffer depth climbs from 4.8 to 8.5 pp at $\delta=0.03$ (Figure 6.4), and on the full-buffer diagnostic from 2.2 to 7.6 pp at $\delta=0.1$ and from 1.5 to 9.1 pp at $\delta=0.3$ (both across all five seeds). Momentum accumulates the already-filtered direction, but the filter attenuates the step, not the source: the new-task loss, and with it the per-step gradient, keeps its full size, and because motion in the braked directions is slow, the shrunk push persists there step after step. A persistent push entering the accumulator is amplified toward $1/(1-\mu)$ times its per-step size, so velocity rebuilds displacement in the very directions the filter shrinks one step at a time. At

the practical buffer this leaves asymmetric PER under momentum at vanilla-level accuracy (0.929) with a depth (8.5 pp) between the scalar gate’s (3.7 pp, C7_M) and vanilla ER’s (15.9 pp). This re-inflation is what makes the head-to-head comparison of the next subsection regime-dependent.

6.4 Reading the Momentum Crosses Together

This subsection answers no separate research question; it reads the momentum crosses of RQ2d and RQ3 side by side, in both gates’ practical regime, the 1k reservoir, with the full-buffer legs staying in their diagnostic role. Together the crosses show that the two gates do not sit on a fixed stability–plasticity frontier that a practitioner trades along: which gate leaves the smaller gap is decided by the momentum setting, the same setting that decides final accuracy. Figure 6.4 shows it directly: the same three methods

at the same buffer, with only the momentum setting changed between the panels.

No momentum: the directional gate is better on both axes. At momentum 0 and a matched 1k buffer, the directional gate beats the scalar gate on both axes (Figure 6.4(a)). Its strongest standard-buffer setting reaches 4.8 pp of depth at 0.880 accuracy (Table 6.2) against the curriculum’s 13.9 pp at 0.837 (C7, Table 6.1): lower gap *and* higher accuracy, because the per-step filter protects the past task without ever slowing the new one at source, so it carries no plasticity debt for the momentum-0 regime to leave unpaid. The one cost that remains is compute: the protection is bought with a conjugate-gradient solve in the replay Fisher at every step, against the scalar gate’s single gradient-norm ratio (Table C.10). When momentum is off and the per-step solve is affordable, the directional gate is the better choice.

Momentum: the scalar gate takes over. Momentum is standard on realistic backbones and, at the training budgets used here, decisive for final accuracy; it reverses the ranking through the re-inflation mechanism of Section 6.3: the directional gate’s depth climbs from 4.8 to 8.5 pp at the matched buffer (Figure 6.4(b)), while the scalar gate composes with momentum because it attenuated at source, and reaches 3.7 pp at vanilla accuracy (C7_M). Once the buffers match, the directional gate has no accuracy advantage left to offset its re-inflated gap (0.929 vs. the curriculum’s 0.931); the scalar gate is simply the lower-gap method in the practically relevant regime. The choice between the gates is therefore not a point on a frontier but a consequence of whether momentum is on: off, take the directional gate; on, take the scalar gate. A five-task rot-MNIST sequence reproduces this picture across four consecutive boundaries: at momentum 0 both gates hold the first task nearly flat through every switch, with the directional gate keeping its edge (the highest final accuracy, 0.870, and the smallest per-boundary area); under momentum the curriculum keeps the least cumulative forgetting, while the re-inflated directional gate loses its edge and ends at vanilla ER’s forgetting level (FORG 0.104 vs. 0.096; Appendix C.9, Figure C.3).

6.5 Generalisation to CIFAR-10 (RQ4)

The CIFAR-10 block tests whether both gates carry to a ResNet-18, a clean $\rightarrow$ corruption shift, and the two normalisations that govern the post-switch recovery. Table 6.3 reports the headline block at momentum 0.9, the practically relevant regime.

Both gates carry to the deeper backbone. The rot-MNIST result carries over and strengthens. Under

batch norm the curriculum cuts depth from 11.8 to 4.8 pp and asymmetric PER to 5.9 pp; under group norm, where vanilla ER suffers a catastrophic 41.2 pp drop, the curriculum cuts it to 5.2 pp and asymmetric PER to 5.6 pp. Every gate-vs-vanilla depth and area contrast holds across all five seeds, at final accuracy within about 0.01 of vanilla and with worst-case retention sharply improved, min-ACC rising from 0.31 to 0.67 under group norm. On this backbone the transient stays open for hundreds of steps, so the area carries as much of the comparison as the depth (Figure 6.5): vanilla ER under group norm crashes T_0 almost to chance and crawls back over the whole recovery window, while both gates hold T_0 near its plateau through the switch. At the matched 5,000 buffer and momentum 0.9 both gates land in the same few-pp band, with the curriculum the cheaper of the two (below), consistent with the momentum regime of Section 6.4. The few-pp dip both gates keep under batch norm is boundary-localised and recovers within tens of steps, consistent with the running statistics re-adapting to the shifted inputs, a pathway no gate on the update touches (Appendix C.6).

Mechanism, generalisation, and cost. Three further checks are reported in full in Appendix C.8. A momentum cross on the deep backbone (Table C.9) turns out to be only weakly informative: at momentum 0 the ResNet under-converges in ten epochs (vanilla ACC 0.65, against 0.730 with momentum), and on that degraded baseline both gates improve with momentum by similar margins (28.6 $\rightarrow$ 5.2 pp for the curriculum, 22.6 $\rightarrow$ 5.6 pp for the directional gate), so the cross cannot separate the predicted composition from a general convergence gain; the vanilla depth null (40.5 $\rightarrow$ 41.2 pp) sits on the same degraded baseline. The momentum signatures therefore rest on the rot-MNIST results, where the MLP isolates them cleanly (Section 7.4). A generalisation block (Appendix C.7, Table C.8) shows the curriculum beating a paired vanilla control on two corruptions it was never tuned on (shot noise 46.3 $\rightarrow$ 7.3 pp, contrast 8.4 $\rightarrow$ 3.9 pp under group norm), with a single three-task group-norm sequence flagged as a reliability boundary of the adaptive ratio when two consecutive corruptions are near-identical and both gradient norms become small (Figure C.2; traced in Appendix C.7). Finally, compute separates the gates where retention does not (Table C.10): the scalar gate is one gradient-norm ratio on top of a standard ER step and leaves the GPU half-idle, while the directional gate’s 25 conjugate-gradient Fisher-vector products per step saturate it, costing $\sim 16\times$ the wall-clock and $\sim 24\times$ the GPU energy on the ResNet-18 at similar peak memory; on the MLP both runs finish in under half a minute and the difference is negligible. In the momentum regime standard in practice, then, the scalar gate matches the directional gate’s retention at




Figure 6.4: The momentum synthesis at the matched 1k buffer: per-step past-task (T_0) accuracy across the rot-MNIST switch for vanilla ER (red), the practical curriculum C7 (blue), and asymmetric PER at its strongest standard-buffer filter $\delta=0.03$ (aqua); seed mean with $\pm$ std bands, $x=0$ the task switch. (a) Momentum 0: the directional gate holds T_0 highest (4.8 pp / ACC 0.880) while the curriculum opens a deep, noisy dip (13.9 pp / 0.837). (b) Momentum 0.9: the ordering swaps: the curriculum holds T_0 near its plateau (3.7 pp / 0.931) while accumulated velocity re-inflates the directional gate’s per-step filter (8.5 pp / 0.929). The pre-switch level is higher in (b) because momentum converges T_0 higher before the switch.

Norm	Method	gap_depth (pp)	gap_area_end	ACC	min-ACC
BatchNorm	G1: Vanilla ER	11.8 ± 2.5	17.5 ± 4.6	0.723 ± 0.008	0.644 ± 0.023
	G2: Curriculum ($\lambda_{\min}=0.20$)	4.8 ± 0.7	3.6 ± 2.2	0.722 ± 0.011	0.713 ± 0.017
	G3: Asym. PER ($\delta=0.1$)	5.9 ± 4.3	3.3 ± 3.0	0.715 ± 0.011	0.698 ± 0.023
GroupNorm	G4: Vanilla ER	41.2 ± 13.5	34.9 ± 23.7	0.730 ± 0.011	0.309 ± 0.111
	G5: Curriculum ($\lambda_{\min}=0.20$)	5.2 ± 2.1	1.7 ± 0.9	0.731 ± 0.020	0.670 ± 0.055
	G6: Asym. PER ($\delta=0.1$)	5.6 ± 1.1	4.8 ± 3.8	0.741 ± 0.009	0.673 ± 0.023

Table 6.3: CIFAR-10 headline block (ResNet-18, momentum 0.9, buffer 5,000, five seeds). Both gates cut gap depth and area far below vanilla ER at matched final accuracy and much higher worst-case retention (min-ACC), under both normalisations. Neither was given a CIFAR-specific sweep; both carry their rot-MNIST configurations.

essentially vanilla-ER cost and without the blind-spot cliffs, which is why it is the default recommendation.

6.6 Summary of Findings

The account holds end to end: the gap separates into a deterministic arc and overshoot of the followed path, closable only by fixing the path and removing the drivers together (RQ1); the scalar feedback gate with momentum reaches 3.7 pp of depth at vanilla accuracy (RQ2); the directional feedforward gate shrinks depth and arc monotonically at no plasticity cost, but fails late and abruptly at its strongest diagnostic setting and re-inflates under momentum (RQ3); the comparison is therefore regime-dependent, and both gates carry to CIFAR-10 at matched accuracy and much improved worst-case retention (RQ4). The discussion reads these findings against the account of Section 3.

7 Discussion

The results support the two-level account and, with it, the gate view of replay. This section reads them back against the theory, weighs the limits of the evidence, and draws out what follows for method design.

7.1 Why the gap needs two levels

The central empirical claim is that the stability gap is not one quantity but two failures of the same optimisation, kept apart because they answer to different interventions; the 2×2 of Section 6.1 is the direct test. Neither intervention is a weaker version of the other. The arc is a property of which path the descent dynamics define, and no amount of optimiser hygiene (exact gradients, balanced magnitudes, zero noise) touches it; the overshoot is how far the discrete, noisy, accelerated iterate leaves whatever path it is on, and no change of path removes it. An inter-




Figure 6.5: Per-step past-task (T_0) accuracy across the CIFAR-10 transition (ResNet-18, momentum 0.9, buffer 5,000), one panel per normalisation, each overlaying vanilla ER (red), the practical curriculum (blue) and asymmetric PER (aqua); seed mean with $\pm$ std bands, $x=0$ the task switch. (a) BatchNorm: vanilla ER dips and re-adapts slowly while both gates hold T_0 near its plateau. (b) GroupNorm: vanilla ER’s transient is deep and long, whereas both gates hold T_0 flat through the switch (Table 6.3, Appendix C.6).

vention aimed at one level is silent on the other.

The natural alternative, a single additive decomposition of the gap into magnitude, variance, and trajectory shares, does not survive the data: the shares are order-dependent (a noisy replay gradient inflates the first-step magnitude), and the mapping from factors to metrics is crossed by the experiments, the curriculum’s momentum-0 tail looking like a “trajectory” term until C8 removes it with the exact buffer and no momentum at all, revealing it as variance (Table C.3). The two-level split is instead by kind: which path is followed versus how tightly it is followed.

The split also settles what the gap is not the price of: C8 closes the transient to 0.2 pp while its accuracy sits 5.5 pp below vanilla ER (Section 6.1), so the permanent rise to the joint-basin level is the stability–plasticity cost and everything below the settled level and back is failure to follow the monotone reference path, not the cost of learning.

Finally, the two levels explain why the metric that matters changes with the backbone. On the MLP the network reconverges in a few steps, so the arc integrates to little and the perceived severity is the overshoot spike; on the ResNet the same transient stays open for hundreds of steps (Figure 6.5), so the slow tail is where most of the lost accuracy accumulates and where both gates’ margin over vanilla is largest (Table 6.3). Depth and area are the signatures of the two levels, and a depth-only reading would have understated the result on exactly the architectures where continual-evaluation cost is highest.

7.2 Two gates, two ways to divide the work

In the unified update of Section 3.3, the curriculum and asymmetric PER occupy the taxonomy’s two pure corners: a scalar gate driven by feedback, and a directional gate driven by feedforward. The results show the two corners working and failing exactly where the taxonomy says they should, and the single fact that organises the whole comparison is *where* the gate attenuates the push.

The scalar feedback gate attenuates at source: it multiplies the entire new-task loss by a schedule that starts near zero, so early in the transition there is almost no push to discretise, accumulate, or let noise act on. Three consequences follow, and the data confirm each. It composes with momentum, because the accumulator never sees a sustained full-strength push to rebuild (Section 6.2). It has no curvature blind spot, because the adaptive schedule reads a live damage signal rather than a model of the past task. And it pays a plasticity debt: holding the loss down slows the new task, a debt only acceleration repays (momentum lifts C8 from 0.818 to 0.955, Table C.3). The scalar gate’s price is plasticity, and momentum is how it is paid.

The directional feedforward gate attenuates per step: it filters the push through the replay Fisher, never slows the new task at source, and so protects the past task at no plasticity cost, ACC and short-horizon plasticity rising as the filter strengthens (Table 6.2). The same per-step locality is its weakness on two axes, one the taxonomy predicts and one that emerged unanticipated. Under momentum the protec-

tion unravels by the re-inflation mechanism of Section 6.3, depth climbing from ≈ 2 to ≈ 8 pp: the same lever that closes the scalar gate’s gap re-opens the directional gate’s, and the difference is entirely where the attenuation is applied, at the source or at the step. That one distinction accounts for the whole momentum column of the study: the vanilla null, the curriculum’s composition, and the directional gate’s re-inflation.

The directional gate’s second failure was not predicted, and it is the sharper one because it makes the case for feedback. At the strongest filter the past-task curve can slide off a late cliff (the $\delta=0.03$ full-buffer cliffs of Section 6.3). A plausible reading is a blind spot of its curvature model: the sampled Fisher rates a direction flat that the true past-task loss punishes further out, and new-task pressure escapes through it. The exposure is characteristic of feedforward control, whose gate can only be as good as its model of the past task; a feedback gate instead reads the damage as it arrives, and its own weakness lies in the conditioning of the signal it reads, as the three-task sequence shows (Section 7.4). The cliffs and the plasticity debt are thus the two gates’ defining and complementary costs. The feedforward gate protects without slowing the new task, but it is only as good as its precomputed model of the past task and can fail where that model is wrong. The feedback gate cannot be fooled that way, because it reads the damage itself, but it pays by holding the whole new task back, a debt only an accelerator repays. This is why the two flip with momentum rather than sitting on a fixed frontier (Section 6.4). Both gates carry to the ResNet at matched accuracy and much improved worst-case retention (Table 6.3), though the momentum cross there is confounded (Section 7.4); it is the gate principle, not either particular method, that generalises.

7.3 What the account changes in the literature

The two-level split settles what reads as a disagreement between the two accounts of the gap that Section 2.2 set out. Lange et al. [16] trace the drop to the gradient imbalance at the boundary; Hess et al. [11] show the gap survives training on the exact joint loss, and Kao et al. [14] derive the geometry that makes it do so. Each account is right about one level and silent on the other. The imbalance is the acute driver of the overshoot: balancing it away is by far the largest single lever on depth in the driver ladder ($19.5 \rightarrow 5.7$ pp), and it leaves the arc untouched. The persistence under exact gradients is the arc: D4 reproduces it under exact, balanced, noiseless replay and measures what remains, 1.9 pp of depth and an area of 3.2, a controlled counterpart of the joint-loss observation and a quantification of the analytical mechanism. Read this way the two accounts stop competing. Most of the

measured depth is the overshoot the imbalance account describes; most of what no repair of the update can remove is the geometry the trajectory account describes; and a mitigation aimed at either alone leaves the other level standing, which is the practical content of the 2×2 .

The unified update of Sodhani et al. [22] sorts replay methods by which weight they move and from what signal; the results add the axis that turned out to be decisive, where the modulation is applied. AGEM and MEGA hold the new-task weight at one and correct per step from a live signal [2, 8], which places them on the step side of the source/step divide, the side whose protection momentum unwinds in Section 6.3. The account therefore predicts that their boundary protection weakens under momentum as the directional gate’s does, while source-attenuating schedules compose with it. These results do not test that prediction, but it is directly testable, and it matters because the original comparisons predate the continual-evaluation lens: end-of-task metrics would not have registered per-step protection eroding, any more than they register the gap itself.

Mode connectivity enters the account as an existence result and exits as a control result. Garipov et al. [6] and, for the continual setting, Mirzadeh et al. [20] establish post hoc that the relevant solutions share a low-loss corridor; the corridor is inspected after training and says nothing about whether an optimiser can be kept inside it while learning. C8 supplies the dynamic counterpart: an update rule whose iterate tracks the monotone reference path through the transition (0.2 pp of depth), turning the observation that a low-loss route exists into the statement that a gate can make training follow it. The same experiment sharpens the stability–plasticity dilemma [19] into a measurement: the dilemma’s price is the permanent offset to the joint level, and it bounds where the model can settle, not how it gets there; the transient that continual evaluation reveals, up to 19.5 pp at the boundary here, is a routing failure the right gate removes.

7.4 Limitations

Several limits bound these claims. The first is statistical: with five seeds the design cannot reach a 0.05 threshold (Section 5.1.2), so the argument rests on effect sizes and seed-paired direction. The effects the story leans on, the up-to-eightfold depth reductions on CIFAR, the C7→C8 variance collapse, the momentum re-inflation of the directional gate, are large relative to the seed spread and unanimous across seeds; the smaller ones, the adjacent- δ steps and the $\lambda_{\min}$ trade-offs, are read as directional consistency only, and one, the adaptive-versus-best-linear depth, does not reach even that ($p = 0.44$).

The second is the accuracy drift of the scalar gate

at momentum 0: lengthening the linear ramp lowers depth but costs several points of final accuracy (Table 6.1), so the method’s accuracy-neutrality is a property of the full curriculum-plus-momentum configuration, not of the ramp alone. The third is the directional gate’s cliffs: the solver control establishes that they are real and not a numerical artefact, but the blind-spot mechanism we offer for them is a reading consistent with the data, not an established cause. The fourth is the CIFAR momentum cross, which we read as uninformative in both directions: at momentum 0 the ResNet under-converges, both gates improve with momentum by similar margins, and the vanilla null sits on the same degraded baseline, so no cell of the deep-backbone cross cleanly confirms or refutes a momentum signature. All three momentum signatures therefore rest on the rot-MNIST results; a clean deep-backbone test would require a converged momentum-0 baseline we did not run. The fifth is the three-task group-norm sequence, where the adaptive ratio loses conditioning at a boundary between two near-identical corruptions and the schedule underperforms its control with high seed variance; we flag it as a reliability boundary of the gradient-ratio feedback signal, absent from the batch-norm variant, rather than a property of the gate principle; Appendix C.7 traces the failure step by step through the logged gradient-norm ratio.

Two scope limits are worth stating plainly. The full-buffer legs of both the driver block and the δ sweep store the entire past task and so break the limited-memory premise of continual learning; they are diagnostic probes that isolate the arc and split momentum’s roles, and the practical configurations are the 1k-buffer curriculum and the standard-buffer directional gate. And the directional gate’s per-step Fisher-vector solve adds a fixed multiple of the backward cost per step, $\sim 16\times$ the wall-clock and $\sim 24\times$ the GPU energy of the curriculum on the ResNet-18 (Table C.10), an overhead the scalar gate’s single gradient-norm ratio does not carry. The headline analysis also rests on a single-transition, two-task design chosen to isolate one boundary; the five-task rot-MNIST sequence (Appendix C.9) shows the regime picture surviving four consecutive boundaries, but a systematic study of longer sequences and of class-incremental rather than domain-incremental shifts is left open.

7.5 Outlook

The taxonomy that organises the two methods also names what is missing from it. The gates studied here fill three of its four corners: scalar $\times$ feedback, directional $\times$ feedforward, and, as an ablation, the linear ramp at scalar $\times$ feedforward. The remaining corner, a directional gate driven by a live damage signal rather than by a precomputed curvature model, is the natural candidate the two failure modes point to: it

would filter the push by direction, keeping the plasticity that makes the feedforward gate cheap, while reading the damage as it arrives rather than trusting a model of the past task. Whether one design can deliver both properties at once is an open question, not an implication of these results, but the momentum results constrain any attempt in advance: its directional attenuation should act at source or on the accumulated velocity, not on the instantaneous step, or it inherits the re-inflation of Section 6.3.

Three narrower directions follow from the failure modes and costs. First, the same repair applies to existing methods: any approach that preconditions the update with past-task curvature per step, the natural-gradient family included [14], carries the re-inflation exposure, and moving its attenuation to the accumulated velocity is a directly testable fix. Second, the directional gate’s compute cost is not fundamental: the per-step conjugate-gradient solve could be replaced by a diagonal or Kronecker-factored Fisher approximation, and the Fisher itself could be refreshed at coarser intervals, up to once per task boundary as Natural Continual Learning does [14]; whether the protection survives the coarser metric or the staler one is open; on backbones much larger than the ResNet-18 here, such an approximation is a precondition for the directional corner rather than a refinement of it. Third, momentum’s gap-closing role proved interchangeable with source variance reduction (C8), so the curriculum could be paired with other variance reducers; iterate averaging, which carries none of momentum’s accumulator-driven overshoot, is the natural candidate.

The cliffs add a methodological rule. They arrive late and abruptly, after long stretches of exemplary retention, and end-of-task metrics would have recorded the $\delta=0.03$ runs as successes; only per-step measurement caught the failure. Continual evaluation [16] is usually presented as the lens that reveals the stability gap. These results argue for it also as an acceptance test for any feedforward gate deployed in a system that serves predictions while it trains: a gate that acts on a model of the past task can fail precisely where that model is wrong, silently until it does.

Statement on the Use of AI Tools

In accordance with the University of Groningen’s guidelines on the responsible use of generative AI, I disclose that AI assistants were used during the preparation of this thesis. The tools were Google’s Gemini 3.x series, Anthropic’s Claude 4.x series, and Anthropic’s Claude Fable 5, the releases current across the period, February to July 2026, in which the work was carried out. They functioned as assistive tools under my continuous direction, not as autonomous contributors. Specifically, I used them to brainstorm and

stress-test ideas and as a sounding board for reasoning through arguments and explanations; to draft and revise prose, which I subsequently edited for accuracy, precision, and voice; to assist in writing and refactoring the analysis and experiment code, all of which I reviewed, ran, and tested; and to produce first-pass summaries of the literature, which I verified against the primary sources.

The tools were never given open-ended control: every request specified a concrete, bounded task, and all output was reviewed and, where retained, corrected, understood, and integrated by me. All experimental results, figures, and numerical claims were checked against the underlying run data, and all cited work was read in the original. The research questions, experimental design, interpretation of the results, and the overall argument of the thesis are my own, and I take full responsibility for its content, including any errors.

References

- [1] Rahaf Aljundi, Min Lin, Baptiste Goujaud, and Yoshua Bengio. Gradient based sample selection for online continual learning, 2019. URL <https://arxiv.org/abs/1903.08671>. 3, 4
- [2] Arslan Chaudhry, Marc’Aurelio Ranzato, Marcus Rohrbach, and Mohamed Elhoseiny. Efficient lifelong learning with a-gem, 2019. URL <https://arxiv.org/abs/1812.00420>. 4, 19
- [3] Arslan Chaudhry, Marcus Rohrbach, Mohamed Elhoseiny, Thalaiyasingam Ajanthan, Puneet K. Dokania, Philip H. S. Torr, and Marc’Aurelio Ranzato. On tiny episodic memories in continual learning, 2019. URL <https://arxiv.org/abs/1902.10486>. 4
- [4] Felix Draxler, Kambis Veschgini, Manfred Salmhofer, and Fred A. Hamprecht. Essentially no barriers in neural network energy landscape, 2019. URL <https://arxiv.org/abs/1803.00885>. 5
- [5] Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis, 2020. URL <https://arxiv.org/abs/1912.05671>. 5
- [6] Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry Vetrov, and Andrew Gordon Wilson. Loss surfaces, mode connectivity, and fast ensembling of dnns, 2018. URL <https://arxiv.org/abs/1802.10026>. 5, 19
- [7] Yiduo Guo, Jie Fu, Huishuai Zhang, Dongyan Zhao, and Yikang Shen. Efficient continual pre-training by mitigating the stability gap, 2024. URL <https://arxiv.org/abs/2406.14833>. 4
- [8] Yunhui Guo, Mingrui Liu, Tianbao Yang, and Tajana Rosing. Improved schemes for episodic memory-based lifelong learning, 2020. URL <https://arxiv.org/abs/1909.11763>. 4, 19
- [9] Md Yousuf Harun and Christopher Kanan. Overcoming the stability gap in continual learning, 2024. URL <https://arxiv.org/abs/2306.01904>. 4
- [10] Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. In *International Conference on Learning Representations (ICLR)*, 2019. URL <https://arxiv.org/abs/1903.12261>. 8, 23
- [11] Timm Hess, Tinne Tuytelaars, and Gido M. van de Ven. Two complementary perspectives to continual learning: Ask not only what to optimize, but also how, 2024. URL <https://arxiv.org/abs/2311.04898>. 3, 4, 19
- [12] Shun ichi Amari. Natural gradient works efficiently in learning. *Neural Computation*, 10(2): 251–276, 1998. 9
- [13] Sandesh Kamath, Albin Soutif-Cormerais, Joost van de Weijer, and Bogdan Raducanu. The expanding scope of the stability gap: Unveiling its presence in joint incremental learning of homogeneous tasks, 2024. URL <https://arxiv.org/abs/2406.05114>. 4
- [14] Ta-Chu Kao, Kristopher T. Jensen, Gido M. van de Ven, Alberto Bernacchia, and Guillaume Hennequin. Natural continual learning: success is a journey, not (just) a destination, 2021. URL <https://arxiv.org/abs/2106.08085>. 3, 4, 6, 19, 20
- [15] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, March 2017. ISSN 1091-6490. doi: 10.1073/pnas.1611835114. URL <http://dx.doi.org/10.1073/pnas.1611835114>. 4
- [16] Matthias De Lange, Gido van de Ven, and Tinne Tuytelaars. Continual evaluation for lifelong learning: Identifying the stability gap, 2023. URL <https://arxiv.org/abs/2205.13452>. 3, 4, 6, 8, 19, 20, 24, 25

- [17] David Lopez-Paz and Marc'Aurelio Ranzato. Gradient episodic memory for continual learning, 2022. URL <https://arxiv.org/abs/1706.08840>. 4
- [18] Michael McCloskey and Neil J. Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *The Psychology of Learning and Motivation*, 24:104–169, 1989. 3
- [19] Martial Mermillod, Aurélie Bugaiska, and Patrick Bonin. The stability-plasticity dilemma: Investigating the continuum from catastrophic forgetting to age-limited learning effects. *Frontiers in psychology*, 4:504, 08 2013. doi: 10.3389/fpsyg.2013.00504. 3, 19
- [20] Seyed Iman Mirzadeh, Mehrdad Farajtabar, Dilan Gorur, Razvan Pascanu, and Hassan Ghasemzadeh. Linear mode connectivity in multitask and continual learning, 2020. URL <https://arxiv.org/abs/2010.04495>. 5, 19
- [21] Roger Ratcliff. Connectionist models of recognition memory: Constraints imposed by learning and forgetting functions. *Psychological Review*, 97:285–308, 04 1990. doi: 10.1037/0033-295X.97.2.285. 3
- [22] Shagun Sodhani, Mojtaba Faramarzi, Saniket Vaibhav Mehta, Pranshu Malviya, Mohamed Abdelsalam, Janarthanan Janarthanan, and Sarath Chandar. An introduction to lifelong supervised learning, 2022. URL <https://arxiv.org/abs/2207.04354>. 3, 4, 19
- [23] Albin Soutif-Cormerais, Antonio Carta, and Joost Van de Weijer. Improving online continual learning performance and stability with temporal ensembles, 2023. URL <https://arxiv.org/abs/2306.16817>. 4
- [24] Edoardo Urettini and Antonio Carta. Online curvature-aware replay: Leveraging 2nd order information for online continual learning, 2025. URL <https://arxiv.org/abs/2502.01866>. 6
- [25] Gido M. van de Ven, Nicholas Soures, and Dhireesha Kudithipudi. *Continual learning and catastrophic forgetting*, page 153–168. Elsevier, 2025. ISBN 9780443157554. doi: 10.1016/b978-0-443-15754-7.00073-0. URL <http://dx.doi.org/10.1016/B978-0-443-15754-7.00073-0>. 4
- [26] Jeffrey S. Vitter. Random sampling with a reservoir. *ACM Trans. Math. Softw.*, 11(1):37–57, March 1985. ISSN 0098-3500. doi: 10.1145/3147.3165. URL <https://doi.org/10.1145/3147.3165>. 4, 23

A Implementation Details

This appendix records the implementation specifics that the methods of Section 5 depend on but that would interrupt the argument in the main text: the datasets, the backbones and training protocol, the experience-replay variants, and the asymmetric-PER solver. The full implementation is available at <https://github.com/Arthurreuss/ValleyWalk>.

A.1 Datasets

rot-MNIST rotates each 28×28 image by the per-task angle and feeds the flattened greyscale vector to the MLP. The CIFAR-10 domain shifts are generated with the `imagecorruptions` package, which implements the standard common-corruptions benchmark of Hendrycks and Dietterich [10]; the Gaussian noise, shot noise, and contrast shifts are applied at severity 3 on the benchmark’s 1–5 scale, per sample at load time, followed by the standard CIFAR-10 channel normalisation. The clean task (“none”) applies the normalisation only.

A.2 Models and Training

Tables A.1 and A.2 give the two backbones, and Table A.3 the training protocol shared across both benchmarks.

Component	Value
Input	784 (flattened 28×28 greyscale)
Hidden layers	2×400 (fully connected, ReLU)
Output	10-class logits
Regularisation	none
Total parameters	478,410

Table A.1: MLP backbone used for the rot-MNIST sweeps.

Component	Value
Input	$3 \times 32 \times 32$ RGB
Stem	3×3 conv, stride 1, no max-pool (CIFAR adaptation)
Stages	4 residual stages, BasicBlocks [2, 2, 2, 2], widths [64, 128, 256, 512]
Normalisation	BatchNorm (default); GroupNorm (min(32, C) groups) variant crossed in Section 5.2.5
Head	global average pool + linear, 10-class logits
Total parameters	$\approx 11.2\text{M}$

Table A.2: CIFAR-adapted ResNet-18 backbone used for the CIFAR-10 block. The 3×3 stride-1 stem and removed max-pool keep the feature map at 32×32 ; all other details follow the ResNet-18.

Hyperparameter	rot-MNIST (MLP)	CIFAR-10 (ResNet-18)
Epochs per task	1 (online)	10
Batch size	256	256
Optimiser	SGD	SGD
Learning rate	0.1	0.1
Momentum	0.0 or 0.9	0.9 (0.0 in cross)
Weight decay	0.0	0.0
Replay buffer	1,000 or 60,000	5,000
Seeds	5	5 or 3

Table A.3: Training protocol for both benchmarks.

A.3 Experience replay

The replay buffer uses reservoir sampling [26], so the stored set is a uniform sample of the past stream at any budget. Three modes appear in the experiments:

- *Standard ER* (the D1 and all C-series conditions): each step takes one combined backward pass on $L_{\text{current}} + L_{\text{replay}}$, with half the mini-batch drawn from the current task and half sampled with replacement from the buffer.

- *Balanced ER* (D3, D4): two separate backward passes whose gradients are individually unit-normalised before being combined 50/50, removing the magnitude asymmetry between g_{new} and g_{replay} .
- *Full-data replay* (D2, D4): the replay gradient is computed on *every* stored sample exactly once per step (no sub-sampling). Paired with a buffer equal to the full past-task training set (60,000 for rot-MNIST T_0), this yields the deterministic empirical past-task gradient at the current parameters, removing estimator variance.

The λ -curriculum re-weights only the current-task term, minimising $L_\lambda = L_{\text{replay}} + \lambda(t) L_{\text{current}}$; the adaptive schedule reads $\|g_{\text{replay}}\|/\|g_{\text{new}}\|$ through an exponential moving average ($\alpha = 0.05$) and applies the optional floor λ_{min} .

A.4 Asymmetric PER

The filtered push $\delta (F_{\text{rep}} + \delta I)^{-1} g_{\text{cur}}$ of Equation (3.4) is computed each step by warm-started conjugate gradients over Fisher-vector products (25 iterations), so the metric is never formed as a matrix; warm-starting from the previous step’s solution keeps the per-step solve cheap across the run. The solver control of Section 5.2.4 reruns the strongest filter ($\delta = 0.03$, full buffer) at 60 iterations to separate properties of the gate from artefacts of an under-converged solve. On the full-buffer leg the replay gradient is exact over all 60k past samples while the Fisher stays estimated on a sampled batch, since a full-buffer Fisher-vector product per CG iteration would dominate the runtime.

B Metric Definitions

This appendix gives the full definitions of every metric recorded by the evaluation pipeline of Section 5.1.2. The main text explains the measures the argument turns on; the tables below are the look-up reference for the complete set.

Stability-gap metrics. Computed from the dense per-step evaluation. Write a_j^{pre} for the pre-task accuracy on task j and $a_j(t)$ for its accuracy at step t of the new task.

Metric	Definition	Role
gap_depth	$\max_j (a_j^{\text{pre}} - \min_t a_j(t))$ over the full record window	Primary metric, worst single point of instability
gap_area_end	$\sum_j \int_0^T \max(0, b_j - a_j(t)) dt$, trapezoidal over the full new-task record (to task end T), with $b_j = \min(a_j^{\text{pre}}, \bar{a}_j^{\text{end}})$ and $\bar{a}_j^{\text{end}}$ the trailing-window mean (last $\max(5, \lceil 0.1n \rceil)$ samples); accuracy at or above b_j contributes 0	Transient cost reported in the results: dip below the recovered level until it closes, excluding permanent forgetting and continued learning (units: accuracy $\cdot$ steps)
recovery_steps	The first step (counted from the transition) at which <i>every</i> past task is simultaneously back at $\geq 90\%$ of its pre-task baseline; undefined (reported as “—”) if no sub-threshold dip occurs, or if recovery is not reached within the record	Duration of the transient, how long the gap stays open

Table B.1: Stability-gap metrics, recorded from dense per-step evaluation. **gap_depth** and **gap_area_end** are the two axes used in the main analysis, how deep the gap cuts and how long it stays open, together with **recovery_steps**.

Continual-learning metrics. Following Lange et al. [16], computed at task boundaries and over the full record. Let $A(E_j, f_n)$ be the accuracy on the test set of task j evaluated at the model f_n after training stage n , and let N be the number of tasks. The full task-boundary accuracy matrix $R[i, j] = A(E_j, f_{T_i})$ is logged for every run, so any downstream metric can be reconstructed.

Metric	Definition	Role
ACC	$\frac{1}{N} \sum_j A(E_j, f_{T_{N-1}})$	Average task accuracy at the final model
FORG	$\frac{1}{N-1} \sum_{i < N-1} [A(E_i, f_{T_i}) - A(E_i, f_{T_{N-1}})]$	Average drop from right-after-learning to final
min-ACC	$\frac{1}{N-1} \sum_{i < N-1} \min\{A(E_i, f_n) : n > T_i\}$	Worst-case retention since learning
WF _w	Mean over tasks of the largest accuracy drop in any window of w consecutive evaluations ($w \in \{10, 100\}$)	Worst-case stability over short / medium horizons
WP _w	Symmetric counterpart of WF _w : the largest accuracy gain ($w \in \{10, 100\}$)	Plasticity / learning velocity
WC-ACC	$\frac{1}{N} A(E_{N-1}, f_{T_{N-1}}) + (1 - \frac{1}{N}) \text{min-ACC}$	Worst-case trade-off, a lower bound on ACC

Table B.2: Continual-learning metrics of Lange et al. [16], recorded at task boundaries and over the full record.

Gradient-fidelity diagnostics. Recorded at the dense cadence for the decomposition conditions only. The tracker compares the replay-buffer gradient g_{replay} against the true past-task gradient g_{true} , computed at every step by a separate forward and backward pass over the entire past-task training set (no sub-sampling, so g_{true} is the deterministic empirical past-task gradient at the current parameters).

Diagnostic	Logged name	Role
$\cos(g_{\text{replay}}, g_{\text{true}})$	<code>true_grad_cosine</code>	Buffer-side directional fidelity, the estimator-variance driver of Section 5.2.1
$\ g_{\text{replay}}\ / \ g_{\text{true}}\ $	<code>true_grad_mag_ratio</code>	Buffer-side magnitude fidelity
$\ g_{\text{new}}\ / \ g_{\text{replay}}\ $	<code>grad_ratio</code>	Update-side magnitude asymmetry, large at step 1 when $\ g_{\text{replay}}\ \approx 0$

Table B.3: Gradient-fidelity diagnostics. Each is logged as a per-step series together with its mean (and minimum, for the cosine). The first two carry the `true_grad` prefix because they are measured against g_{true} ; `grad_ratio` compares g_{replay} to the current-task gradient g_{new} and its inverse drives the adaptive curriculum of Section 5.2.2.

C Complete Results

This appendix gives the full metric set for every condition of the main experimental blocks, as recorded by the aggregation pipeline of Section 5.1.2. Each row is the mean over five seeds (three for the CIFAR generalisation block) with the seed standard deviation. Accuracy-type metrics (ACC, FORG, min-ACC, WF₁₀₀, WP₁₀₀, WC-ACC) are fractions; `gap_depth` is in percentage points (pp); `gap_area_end` is the transient cost (accuracy-fraction $\times$ steps, Table B.1), the dip below the level each task recovers to, so permanent forgetting does not contribute; “Recov.” is `recovery_steps`, with “—” marking conditions where every past task stayed above the 90% recovery threshold for the whole window (no recovery event to time).

C.1 rot-MNIST: Decomposition (D-series)

Condition	ACC	FORG	min-ACC	WF ₁₀₀	WP ₁₀₀	WC-ACC	Depth	Area end	Recov.
<i>Momentum 0.0</i>									
D1 vanilla	0.873 $\pm$ 0.005	0.019 $\pm$ 0.020	0.706 $\pm$ 0.056	0.147 $\pm$ 0.029	0.736 $\pm$ 0.020	0.794 $\pm$ 0.028	19.5 $\pm$ 4.1	6.3 $\pm$ 0.5	3.6 $\pm$ 0.9
D2 fulldata	0.890 $\pm$ 0.009	-0.016 $\pm$ 0.009	0.719 $\pm$ 0.039	0.126 $\pm$ 0.020	0.736 $\pm$ 0.018	0.801 $\pm$ 0.022	18.3 $\pm$ 2.5	8.1 $\pm$ 0.7	3.0 $\pm$ 1.2
D3 balanced	0.849 $\pm$ 0.005	0.031 $\pm$ 0.018	0.824 $\pm$ 0.006	0.045 $\pm$ 0.002	0.721 $\pm$ 0.004	0.836 $\pm$ 0.004	5.7 $\pm$ 1.9	1.4 $\pm$ 0.4	—
D4 fulldata balanced	0.869 $\pm$ 0.002	-0.014 $\pm$ 0.013	0.862 $\pm$ 0.002	0.030 $\pm$ 0.007	0.710 $\pm$ 0.003	0.852 $\pm$ 0.002	1.9 $\pm$ 1.3	3.2 $\pm$ 0.5	—
<i>Momentum 0.9</i>									
D1 vanilla	0.928 $\pm$ 0.003	0.058 $\pm$ 0.007	0.797 $\pm$ 0.028	0.092 $\pm$ 0.016	0.808 $\pm$ 0.012	0.878 $\pm$ 0.014	15.9 $\pm$ 2.6	1.0 $\pm$ 0.2	12.6 $\pm$ 1.7
D2 fulldata	0.964 $\pm$ 0.002	-0.010 $\pm$ 0.003	0.807 $\pm$ 0.009	0.082 $\pm$ 0.004	0.808 $\pm$ 0.012	0.885 $\pm$ 0.004	14.9 $\pm$ 0.9	3.1 $\pm$ 0.2	10.6 $\pm$ 1.1
D3 balanced	0.932 $\pm$ 0.005	0.038 $\pm$ 0.003	0.897 $\pm$ 0.006	0.040 $\pm$ 0.003	0.822 $\pm$ 0.011	0.922 $\pm$ 0.003	5.9 $\pm$ 0.6	0.4 $\pm$ 0.2	—
D4 fulldata balanced	0.967 $\pm$ 0.001	-0.022 $\pm$ 0.003	0.942 $\pm$ 0.005	0.016 $\pm$ 0.004	0.834 $\pm$ 0.010	0.950 $\pm$ 0.002	1.3 $\pm$ 0.4	0.2 $\pm$ 0.1	—

Table C.1: rot-MNIST decomposition block, full metrics, at both momentum settings. Buffer-fidelity diagnostics for these runs are in Table C.4.




Figure C.1: The full driver ladder: per-step past-task (T_0) accuracy across the $T_0 \rightarrow T_1$ transition for the four driver-removal conditions at momentum 0; solid lines are the seed mean, bands the standard deviation, $x=0$ the task switch. The boundary spike falls when balancing switches off the imbalance driver (D1, D2 vs. D3, D4); the recovery is smoothest where the imbalance is off as well (D4, and to a lesser extent D3), while D2 keeps a visible jitter from its unbalanced, still-stochastic new-task push despite the exact replay gradient. The dip-and-recover bend that no removal touches, cleanest in D4, is the deterministic arc (Section 6.1).

C.2 rot-MNIST: λ -Curriculum (C-series)

Condition	ACC	FORG	min-ACC	WF ₁₀₀	WP ₁₀₀	WC-ACC	Depth	Area end	Recov.
<i>Momentum 0.0</i>									
C1 linear N50	0.851 $\pm$ 0.023	0.042 $\pm$ 0.021	0.707 $\pm$ 0.019	0.148 $\pm$ 0.011	0.748 $\pm$ 0.006	0.784 $\pm$ 0.021	17.5 $\pm$ 1.7	4.9 $\pm$ 0.5	35.8 $\pm$ 6.3
C2 linear N100	0.844 $\pm$ 0.025	0.046 $\pm$ 0.023	0.719 $\pm$ 0.033	0.140 $\pm$ 0.025	0.730 $\pm$ 0.006	0.786 $\pm$ 0.019	16.2 $\pm$ 4.3	3.6 $\pm$ 0.3	64.0 $\pm$ 8.6
C3 linear N200	0.828 $\pm$ 0.025	0.054 $\pm$ 0.024	0.739 $\pm$ 0.048	0.128 $\pm$ 0.043	0.688 $\pm$ 0.004	0.784 $\pm$ 0.034	14.3 $\pm$ 4.6	1.8 $\pm$ 0.6	132.8 $\pm$ 43.4
C4 adaptive	0.833 $\pm$ 0.025	0.050 $\pm$ 0.019	0.754 $\pm$ 0.025	0.111 $\pm$ 0.016	0.703 $\pm$ 0.007	0.794 $\pm$ 0.020	12.8 $\pm$ 2.9	2.2 $\pm$ 0.5	88.8 $\pm$ 6.1
C5 adaptive lmin0.05	0.835 $\pm$ 0.025	0.048 $\pm$ 0.019	0.749 $\pm$ 0.027	0.113 $\pm$ 0.016	0.702 $\pm$ 0.006	0.793 $\pm$ 0.020	13.2 $\pm$ 3.2	2.2 $\pm$ 0.5	88.8 $\pm$ 6.1
C6 adaptive lmin0.10	0.836 $\pm$ 0.026	0.048 $\pm$ 0.021	0.743 $\pm$ 0.026	0.120 $\pm$ 0.019	0.700 $\pm$ 0.007	0.791 $\pm$ 0.019	13.9 $\pm$ 3.3	2.4 $\pm$ 0.5	88.4 $\pm$ 11.0
C7 adaptive lmin0.20	0.837 $\pm$ 0.031	0.049 $\pm$ 0.024	0.743 $\pm$ 0.030	0.119 $\pm$ 0.023	0.710 $\pm$ 0.010	0.792 $\pm$ 0.024	13.9 $\pm$ 3.4	2.7 $\pm$ 0.6	64.8 $\pm$ 8.2
<i>Momentum 0.9</i>									
C1 linear N50	0.932 $\pm$ 0.003	0.053 $\pm$ 0.005	0.892 $\pm$ 0.004	0.041 $\pm$ 0.004	0.830 $\pm$ 0.012	0.927 $\pm$ 0.003	6.4 $\pm$ 0.5	0.2 $\pm$ 0.1	—
C2 linear N100	0.932 $\pm$ 0.003	0.051 $\pm$ 0.008	0.892 $\pm$ 0.005	0.036 $\pm$ 0.003	0.826 $\pm$ 0.012	0.926 $\pm$ 0.002	6.4 $\pm$ 0.6	0.2 $\pm$ 0.1	—
C3 linear N200	0.925 $\pm$ 0.007	0.051 $\pm$ 0.006	0.893 $\pm$ 0.004	0.031 $\pm$ 0.004	0.819 $\pm$ 0.010	0.920 $\pm$ 0.006	6.2 $\pm$ 0.6	0.1 $\pm$ 0.0	—
C4 adaptive	0.917 $\pm$ 0.001	0.019 $\pm$ 0.004	0.931 $\pm$ 0.005	0.016 $\pm$ 0.003	0.798 $\pm$ 0.010	0.915 $\pm$ 0.001	2.5 $\pm$ 0.5	0.3 $\pm$ 0.1	—
C5 adaptive lmin0.05	0.922 $\pm$ 0.001	0.023 $\pm$ 0.004	0.930 $\pm$ 0.005	0.016 $\pm$ 0.002	0.802 $\pm$ 0.009	0.921 $\pm$ 0.001	2.6 $\pm$ 0.5	0.1 $\pm$ 0.0	—
C6 adaptive lmin0.10	0.927 $\pm$ 0.002	0.028 $\pm$ 0.005	0.926 $\pm$ 0.004	0.017 $\pm$ 0.003	0.806 $\pm$ 0.009	0.927 $\pm$ 0.002	3.0 $\pm$ 0.4	0.0 $\pm$ 0.0	—
C7 adaptive lmin0.20	0.931 $\pm$ 0.002	0.035 $\pm$ 0.005	0.919 $\pm$ 0.003	0.021 $\pm$ 0.004	0.812 $\pm$ 0.009	0.930 $\pm$ 0.001	3.7 $\pm$ 0.4	0.0 $\pm$ 0.0	—

Table C.2: rot-MNIST λ -curriculum sweep, full metrics, at both momentum settings. All conditions applied on top of standard ER (1 k reservoir).

Full-buffer momentum diagnostic (C8). Table C.3 crosses the practical curriculum with the full 60k buffer at both momentum settings, isolating momentum’s two roles (Section 6.2). Source variance reduction alone (C8, full buffer, $\mu=0$) already closes the gap to 0.2 pp, so the residual the curriculum leaves at $\mu=0$ is estimator variance, not structure; momentum’s gap-closing role is interchangeable with the full buffer. Acceleration alone (the $\mu=0.9$ column) restores accuracy: C8_M reaches 0.955, approaching the full-data ceiling (D2_M: 0.964), while re-introducing a small transient (2.0 pp, area 0.38). The shape and timing of this transient point to the $\lambda_{\min}$ floor meeting momentum. With the replay gradient exact and the arc removed, the only push left in the first steps is the floored fifth of the new-task gradient (the logged gradient-norm ratio confirms the schedule sits at $\lambda = 0.2$ through the dip window), and momentum amplifies a persistent push toward $1/(1 - \mu) = 10$ times its per-step size: the dip accordingly opens at the boundary, bottoms out about ten steps in, and closes within about fifty as the replay gradient reasserts itself. The floor is by construction the one share of the push the source gate does not attenuate, so momentum briefly re-inflates exactly that share, a miniature of the directional gate’s re-inflation. The dose-response on the same exact buffer supports the reading: the floored push without momentum costs 0.2 pp (C8), with momentum 2.0 pp (C8_M), and the unfloored abrupt switch with momentum 14.9 pp (D2_M).

$C8_M$'s area additionally counts in full under the end-referencing convention of Section 5.1.2, since it recovers above its pre-switch level. $C8/C8_M$ are diagnostic probes: the full buffer defeats the limited-memory premise, and the practical scalar-gate configuration remains $C7_M$ (3.7 pp / 0.03 / 0.931) at 1/60th the memory.

	C7 (1k, $\mu=0$)	$C7_M$ (1k, $\mu=0.9$)	C8 (full, $\mu=0$)	$C8_M$ (full, $\mu=0.9$)
gap_depth (pp)	13.9	3.7	0.2	2.0
gap_area_end	2.7	0.03	0.05	0.38
ACC	0.837	0.931	0.818	0.955

Table C.3: Curriculum $\times$ full buffer $\times$ momentum on rot-MNIST, isolating momentum's two roles. Source variance reduction alone (C8) closes the gap; acceleration alone (the μ column) restores accuracy. The residual at C7 ($\mu=0$) is variance, not structure.

C.3 rot-MNIST: Buffer-Fidelity Diagnostics

The gradient diagnostics of Appendix B were logged for the decomposition block only. The full-data conditions (D2, D4) reach exact fidelity by construction, the replay gradient equals the deterministic empirical past-task gradient, so cosine and magnitude ratio are 1.000. The finite-buffer conditions (D1, D3) show the directional and magnitude error a 1,000-sample reservoir incurs against g_{true} .

Condition	cōs	cos _{min}	$\bar{\rho}$
D1 vanilla ($\mu=0$)	0.711 ± 0.035	0.086 ± 0.103	1.040 ± 0.039
D3 balanced ($\mu=0$)	0.665 ± 0.043	0.025 ± 0.140	1.137 ± 0.060
D1 vanilla ($\mu=0.9$)	0.551 ± 0.033	0.135 ± 0.069	0.418 ± 0.014
D3 balanced ($\mu=0.9$)	0.441 ± 0.039	-0.052 ± 0.117	0.265 ± 0.024
D2, D4 (full data)	1.000	1.000	1.000

Table C.4: Buffer-fidelity diagnostics for the decomposition block: cōs and cos_{min} are the mean and minimum of $\cos(g_{\text{replay}}, g_{\text{true}})$ over the window; $\bar{\rho}$ is the mean of $\|g_{\text{replay}}\|/\|g_{\text{true}}\|$.

C.4 rot-MNIST: Asymmetric PER (P-series)

Table C.5 gives the full metric set for the asymmetric-PER δ sweep on both buffer legs and at both momentum settings; the momentum-0 depth/area/ACC columns match Table 6.2 in the main text, and the momentum-0.9 rows are the source of the re-inflation numbers of Section 6.3. Momentum-0.9 runs cover the three strongest filters on the standard leg and $\delta \in \{0.3, 0.1\}$ on the full-buffer leg.

Condition	ACC	FORG	min-ACC	WF ₁₀₀	WP ₁₀₀	WC-ACC	Depth	Area end
<i>Standard (1k) buffer, momentum 0.0</i>								
$\delta=1.0$	0.869 ± 0.005	0.027 ± 0.018	0.792 ± 0.012	0.081 ± 0.009	0.733 ± 0.009	0.838 ± 0.006	9.0 ± 0.6	3.29 ± 0.65
$\delta=0.3$	0.873 ± 0.004	0.023 ± 0.017	0.807 ± 0.012	0.068 ± 0.009	0.736 ± 0.008	0.847 ± 0.005	7.4 ± 1.1	2.32 ± 0.59
$\delta=0.1$	0.876 ± 0.004	0.018 ± 0.016	0.822 ± 0.012	0.056 ± 0.009	0.739 ± 0.008	0.856 ± 0.005	5.9 ± 1.5	1.56 ± 0.65
$\delta=0.03$	0.880 ± 0.005	0.012 ± 0.015	0.834 ± 0.012	0.048 ± 0.005	0.741 ± 0.007	0.862 ± 0.005	4.8 ± 1.9	1.04 ± 0.58
<i>Standard (1k) buffer, momentum 0.9</i>								
$\delta=0.3$	0.929 ± 0.002	0.061 ± 0.004	0.851 ± 0.008	0.065 ± 0.004	0.813 ± 0.009	0.907 ± 0.005	10.4 ± 0.5	0.50 ± 0.17
$\delta=0.1$	0.928 ± 0.002	0.061 ± 0.006	0.861 ± 0.008	0.060 ± 0.005	0.815 ± 0.009	0.911 ± 0.005	9.4 ± 0.5	0.38 ± 0.19
$\delta=0.03$	0.929 ± 0.003	0.060 ± 0.009	0.871 ± 0.008	0.053 ± 0.007	0.816 ± 0.010	0.917 ± 0.005	8.5 ± 0.6	0.23 ± 0.12
<i>Full (60k) buffer, momentum 0.0</i>								
$\delta=1.0$	0.898 ± 0.002	-0.025 ± 0.012	0.841 ± 0.011	0.037 ± 0.011	0.734 ± 0.006	0.865 ± 0.006	4.0 ± 1.0	2.32 ± 1.38
$\delta=0.3$	0.900 ± 0.002	-0.029 ± 0.012	0.867 ± 0.002	0.026 ± 0.007	0.739 ± 0.005	0.878 ± 0.002	1.5 ± 1.1	1.01 ± 0.91
$\delta=0.1$	0.901 ± 0.002	-0.032 ± 0.013	0.860 ± 0.026	0.040 ± 0.028	0.741 ± 0.004	0.875 ± 0.013	2.2 ± 1.6	0.44 ± 0.38
$\delta=0.03$	0.903 ± 0.002	-0.035 ± 0.012	0.787 ± 0.094	0.118 ± 0.102	0.742 ± 0.004	0.838 ± 0.048	9.5 ± 8.2	0.43 ± 0.25
$\delta=0.03$ (CG 60)	0.903 ± 0.002	-0.035 ± 0.012	0.766 ± 0.113	0.126 ± 0.109	0.742 ± 0.004	0.828 ± 0.057	11.5 ± 10.1	0.47 ± 0.28
<i>Full (60k) buffer, momentum 0.9</i>								
$\delta=0.3$	0.964 ± 0.001	-0.012 ± 0.002	0.864 ± 0.009	0.052 ± 0.005	0.813 ± 0.008	0.913 ± 0.005	9.1 ± 0.8	1.91 ± 0.26
$\delta=0.1$	0.965 ± 0.001	-0.012 ± 0.003	0.880 ± 0.008	0.044 ± 0.004	0.814 ± 0.008	0.921 ± 0.005	7.6 ± 0.7	1.57 ± 0.25

Table C.5: rot-MNIST asymmetric-PER δ sweep, full metrics: both buffer legs at both momentum settings, plus the solver control (CG 60). The momentum-0.9 rows show the re-inflation of Section 6.3: depth climbs on every leg relative to its momentum-0 counterpart, while the full-buffer accuracies (≈ 0.965) are diagnostic values only, the full buffer defeating the limited-memory premise.

C.5 rot-MNIST: Symmetric vs. Asymmetric Preconditioning

Leaving the restoring gradient g_{rep} at full strength is the directional gate's defining choice (Section 6.3). The alternative is to precondition the summed update, $d = \delta(F_{\text{rep}} + \delta I)^{-1}(g_{\text{cur}} + g_{\text{rep}})$, the symmetric variant that

passes both gradients through the same metric. Because g_{rep} is the gradient of the replay loss, it lies in the top eigenspace of that loss’s own Fisher F_{rep} , so the symmetric filter shrinks the restoring gradient *harder* than the harmful push it is meant to brake and cannot preferentially protect the past task. Table C.6 sweeps the same δ on both variants, and they are mirror images. As the filter strengthens ($\delta \rightarrow 0$) the asymmetric gate’s forgetting falls (FORG 0.027 $\rightarrow$ 0.012) and its accuracy rises, while the symmetric gate’s forgetting *climbs* (0.030 $\rightarrow$ 0.066) and its accuracy falls; every one of these four trends holds across all five seeds. Symmetric depth barely moves across the sweep (5.7 $\rightarrow$ 6.7 pp) and at $\delta=1.0$ even sits below the asymmetric gate’s (5.7 vs. 9.0 pp), because filtering the whole update also damps the per-step *expression* of the dip. This is the area caveat of Section 5.1.2 made concrete: the symmetric gate’s end-referenced area collapses toward zero (2.7 $\rightarrow$ 0.01) not because the dip recovers but because it stops recovering, settling into a permanent accommodation that the area excludes and the forgetting metric records. Read on depth or transient area the symmetric gate looks competitive; read on forgetting and final accuracy it degrades monotonically as its filter strengthens. Asymmetry is what turns the filter from a reshuffling of the damage into a net removal of it.

δ	Asymmetric (g_{rep} free)				Symmetric (joint)			
	depth (pp)	area	FORG	ACC	depth (pp)	area	FORG	ACC
1.0	9.0	3.29	0.027	0.869	5.7	2.69	0.030	0.872
0.3	7.4	2.32	0.023	0.873	5.9	1.65	0.042	0.867
0.1	5.9	1.56	0.018	0.876	6.2	0.34	0.057	0.861
0.03	4.8	1.04	0.012	0.880	6.7	0.01	0.066	0.857

Table C.6: Asymmetric versus symmetric preconditioning on the standard 1k buffer at momentum 0, five seeds, sweeping the shared damping δ . Both variants filter through the replay Fisher F_{rep} ; they differ only in whether the metric is applied to the new-task push alone (asymmetric, g_{rep} left free) or to the summed update (symmetric). As $\delta \rightarrow 0$ the asymmetric gate improves on every axis while the symmetric gate degrades on every axis; symmetric depth stays flat because the filter suppresses the per-step expression of the dip even as the underlying forgetting grows, and the symmetric end-referenced area collapses toward zero for the same reason: the past task stops recovering, so the loss registers in FORG rather than in the transient area (Section 5.1.2).

C.6 CIFAR-10: Headline Block

Condition	ACC	FORG	min-ACC	WF ₁₀₀	WP ₁₀₀	WC-ACC	Depth	Area end
G1 vanilla BN	0.723 $\pm$ 0.008	0.010 $\pm$ 0.019	0.644 $\pm$ 0.023	0.127 $\pm$ 0.015	0.484 $\pm$ 0.020	0.667 $\pm$ 0.016	11.8 $\pm$ 2.5	17.5 $\pm$ 4.6
G2 adaptive BN	0.722 $\pm$ 0.011	0.006 $\pm$ 0.016	0.713 $\pm$ 0.017	0.070 $\pm$ 0.010	0.469 $\pm$ 0.016	0.698 $\pm$ 0.013	4.8 $\pm$ 0.7	3.6 $\pm$ 2.2
G3 asym. PER BN	0.715 $\pm$ 0.011	-0.008 $\pm$ 0.018	0.698 $\pm$ 0.023	0.124 $\pm$ 0.022	0.460 $\pm$ 0.025	0.682 $\pm$ 0.023	5.9 $\pm$ 4.3	3.3 $\pm$ 3.0
G4 vanilla GN	0.730 $\pm$ 0.011	-0.023 $\pm$ 0.030	0.309 $\pm$ 0.111	0.294 $\pm$ 0.097	0.415 $\pm$ 0.058	0.507 $\pm$ 0.054	41.2 $\pm$ 13.5	34.9 $\pm$ 23.7
G5 adaptive GN	0.731 $\pm$ 0.020	-0.026 $\pm$ 0.026	0.670 $\pm$ 0.055	0.095 $\pm$ 0.033	0.340 $\pm$ 0.052	0.687 $\pm$ 0.037	5.2 $\pm$ 2.1	1.7 $\pm$ 0.9
G6 asym. PER GN	0.741 $\pm$ 0.009	-0.035 $\pm$ 0.018	0.673 $\pm$ 0.023	0.080 $\pm$ 0.009	0.330 $\pm$ 0.024	0.696 $\pm$ 0.015	5.6 $\pm$ 1.1	4.8 $\pm$ 3.8

Table C.7: CIFAR-10 headline block, full metrics (ResNet-18, momentum 0.9, buffer 5,000, five seeds). “adaptive” is the practical curriculum with $\lambda_{\text{min}} = 0.20$; “asym. PER” runs at $\delta = 0.1$.

The gates’ residual batch-norm dip. The small gap the gates keep under batch norm is a genuine boundary transient: in Figure 6.5(a) both gated curves dip together over the first steps after the switch, by roughly the tabulated 5 pp, and recover to their plateau within a few tens of steps. A reading consistent with this pattern is the normalisation pathway of Section 5.2.5. BatchNorm’s running statistics update on every forward pass, whatever the gate does to the gradient, so the first corrupted batches shift them away from the clean task and no gate on the update can hold that back; the clean replay samples in every batch pull the statistics back, which matches the fast recovery. The group-norm contrast supports the reading: with no cross-batch statistics to disturb, the curriculum holds T_0 flat through the same boundary (Figure 6.5(b)) and its transient area falls to roughly half its batch-norm value (1.7 vs. 3.6). This is a reading rather than an established mechanism; a control that freezes the running statistics across the switch would isolate it.

C.7 CIFAR-10: Generalisation Block

Each shift pairs the practical curriculum (“curric.”: adaptive $\lambda_{\text{min}} = 0.20$, buffer 5,000) against a vanilla-ER control on the same shift, at the indicated normalisation, over three seeds.

Condition	ACC	FORG	min-ACC	WC-ACC	Depth	Area end
<i>BatchNorm</i>						
shot noise (vanilla)	0.724 ± 0.011	0.017 ± 0.010	0.633 ± 0.017	0.663 ± 0.014	13.1 ± 3.0	21.2 ± 3.9
shot noise (curric.)	0.723 ± 0.014	0.010 ± 0.010	0.708 ± 0.032	0.697 ± 0.016	5.6 ± 1.0	5.8 ± 1.5
contrast (vanilla)	0.671 ± 0.026	-0.008 ± 0.006	0.610 ± 0.063	0.586 ± 0.048	15.4 ± 3.9	33.4 ± 4.5
contrast (curric.)	0.671 ± 0.044	-0.012 ± 0.007	0.727 ± 0.034	0.643 ± 0.048	3.7 ± 1.1	2.2 ± 0.9
3-task (vanilla)	0.740 ± 0.006	-0.011 ± 0.013	0.670 ± 0.015	0.691 ± 0.010	3.3 ± 0.4	6.0 ± 0.7
3-task (curric.)	0.745 ± 0.010	-0.021 ± 0.008	0.677 ± 0.025	0.696 ± 0.017	6.9 ± 2.7	3.9 ± 0.1
<i>GroupNorm</i>						
shot noise (vanilla)	0.718 ± 0.021	-0.036 ± 0.036	0.244 ± 0.075	0.469 ± 0.032	46.3 ± 12.0	54.4 ± 40.8
shot noise (curric.)	0.705 ± 0.061	-0.028 ± 0.012	0.638 ± 0.108	0.656 ± 0.087	7.3 ± 5.9	2.4 ± 1.1
contrast (vanilla)	0.777 ± 0.022	-0.067 ± 0.021	0.641 ± 0.061	0.710 ± 0.040	8.4 ± 4.4	3.3 ± 0.6
contrast (curric.)	0.778 ± 0.032	-0.070 ± 0.013	0.668 ± 0.077	0.724 ± 0.054	3.9 ± 3.1	1.2 ± 0.8
3-task (vanilla)	0.722 ± 0.020	-0.023 ± 0.015	0.468 ± 0.043	0.550 ± 0.023	4.4 ± 0.4	10.3 ± 3.6
3-task (curric.)	0.721 ± 0.019	-0.025 ± 0.020	0.421 ± 0.114	0.517 ± 0.070	32.8 ± 14.5	32.0 ± 19.5

Table C.8: CIFAR-10 generalisation block, full metrics. The two-task novel corruptions (shot noise, contrast) favour the curriculum under both normalisations; the three-task group-norm sequence is the outlier analysed below.

The three-task group-norm failure. The three-task sequence is the one place the curriculum underperforms its control (depth 32.8 ± 14.5 pp against vanilla’s 4.4 ± 0.4 , Table C.8; Figure C.2), and the logged gradient-norm traces locate the failure in the conditioning of the feedback signal. The adaptive ratio $r = \|g_{\text{replay}}\|/\|g_{\text{new}}\|$ carries information through the boundary ordering of Section 5.2.2: at a well-separated switch the new-task gradient dwarfs the replay gradient, so r starts near zero and its rise tracks the replay gradient reasserting itself. The first boundary of this sequence behaves that way (clean $\rightarrow$ Gaussian: first-step new-task loss 1.9–2.1, $r \approx 0.2$), and the schedule holds every past task flat through it. The second boundary never establishes the ordering: Gaussian and shot noise are near-identical, so the incoming model already fits the new task (T_2 accuracy 0.67–0.71 before its first step, new-task loss 0.5–0.7), leaving $\|g_{\text{new}}\|$ small, while $\|g_{\text{replay}}\|$ is small because the model is converged on the replayed tasks. The ratio becomes a quotient of two near-zero norms: under group norm it reads 0.5–1.3 from the first step and fluctuates step to step, so the EMA drives λ from the floor to ≈ 0.7 –0.8 within fifty steps and back to $\lambda_{\min}$ within a few hundred. The gate is moved by noise, not by damage.

What follows, in all three seeds, is not a stability gap but an optimisation instability. Within the first forty steps the replay loss itself diverges (0.2 – $0.4 \rightarrow 0.8$ – 2.1) and every task falls together, *including the one being trained* (T_2 : $0.71 \rightarrow 0.29$, $0.69 \rightarrow 0.30$, $0.67 \rightarrow 0.51$ across seeds), before the run recovers; the per-seed depth scales with the divergence, which is where the ± 14.5 pp spread comes from. Because all tasks crash together, the depth and area of the three-task group-norm rows record this instability, not the past-task transient the rest of the study measures. Two controls isolate the trigger. The vanilla run crosses the same boundary with a flat replay loss and 4.4 pp of depth, so the near-identical boundary is harmless on its own; the batch-norm run executes the identical schedule through the same boundary, with the same ill-conditioned early ratio, without instability (Figure C.2), so the noise-driven λ is not sufficient either. The failure needs both the unconditioned signal and the group-norm backbone’s boundary sensitivity, consistent with the far deeper group-norm transients vanilla ER shows on the same corruption family (shot noise: 46.3 vs 13.1 pp, Table C.8). We therefore read the episode as a reliability boundary of the gradient-norm-ratio signal rather than of the gate principle: the ratio loses its meaning when both norms become small, and the design admits direct guards: freezing λ when both norms fall below a reference scale, or reading a damage signal that does not vanish at a quiet boundary, such as the replay-loss excess over its pre-switch level.




Figure C.2: Per-step past-task accuracy across the three-task CIFAR-10 generalisation sequence (none $\rightarrow$ Gaussian $\rightarrow$ shot, ResNet-18, practical curriculum, three seeds with $\pm$ std bands), comparing the group-norm run (orange) against its batch-norm counterpart (blue). The two switches (dashed) fall at $\sim 2k$ and $\sim 4k$ steps. The first boundary (none $\rightarrow$ Gaussian) barely perturbs the past tasks in either norm. The second (Gaussian $\rightarrow$ shot, two near-identical corruptions) drives a deep transient in T_0 (a) and T_1 (b) under group norm only, while the batch-norm run holds flat through both switches. The dip is thus localised entirely to the second transition and to group norm, the adaptive-ratio reliability boundary analysed above.

C.8 CIFAR-10: Momentum Cross and Compute Cost

This subsection collects the three deep-backbone checks summarised in Section 6.5: the momentum cross (Table C.9), the generalisation block (Table C.8 above), and the compute cost (Table C.10).

Momentum cross. The cross repeats the composability test of Section 5.2.3 on the ResNet, pairing momentum-0 runs at group norm against the momentum-0.9 rows, and it turns out to be only weakly informative. At momentum 0 the ResNet under-converges the past task in ten epochs (vanilla min-ACC 0.11, ACC 0.65), so every momentum-0 depth sits on a degraded baseline. On that baseline both gates improve with momentum by similar margins (curriculum $28.6 \rightarrow 5.2$ pp, asymmetric PER $22.6 \rightarrow 5.6$ pp), so the cross cannot separate the predicted source-attenuation composition from a general convergence gain, and the prediction that momentum re-inflates the directional gate does not appear; the vanilla null ($40.5 \rightarrow 41.2$ pp) is directionally consistent with the account but rests on the same degraded baseline. We therefore draw no mechanism conclusion from this cross in either direction: the momentum signatures rest on the rot-MNIST results, where the MLP isolates them cleanly (Sections 6.2 and 6.3), and a clean deep-backbone test would need a converged momentum-0 baseline (Section 7.4).

Method (GroupNorm)	depth, $\mu=0$	depth, $\mu=0.9$
Vanilla ER	40.5 ± 4.0	41.2 ± 13.5
Curriculum ($\lambda_{\min}=0.20$)	28.6 ± 8.8	5.2 ± 2.1
Asym. PER ($\delta=0.1$)	22.6 ± 9.7	5.6 ± 1.1

Table C.9: CIFAR-10 momentum cross (group norm, buffer 5,000, five seeds), gap depth in pp. All momentum-0 cells sit on an under-converged baseline (vanilla min-ACC 0.11), so the cross is read as uninformative for the momentum signatures rather than as partial confirmation (see text).

Compute cost. Cost is measured with a dedicated benchmark: one full two-task run per gate in its headline configuration (curriculum: adaptive schedule, $\lambda_{\min}=0.20$; asymmetric PER: $\delta=0.1$, 25 warm-started CG iterations; momentum 0.9, group norm and buffer 5,000 on CIFAR), five seeds each on a single NVIDIA RTX

PRO 6000, with per-step stability evaluation disabled so the numbers reflect method compute alone. Wall-clock and CPU time come from `getrusage`; GPU utilisation, power, and memory are sampled at 1 Hz with `nvidia-smi`, and energy is the integrated power. On the ResNet-18 the split is stark (Table C.10). The scalar gate adds one gradient-norm ratio and an EMA update to a standard ER step, and its run is not even compute-bound: the GPU sits half-idle at 49% utilisation and 259 W. The directional gate’s 25 conjugate-gradient Fisher-vector products per step, each on the order of a backward pass, saturate the same card (96%, 367 W) for $16\times$ as long, and since energy is power times time the two factors compound to a $24\times$ energy gap (13 vs. 320 Wh per run). Memory is not the price: peak GPU memory differs by 11%. On the rot-MNIST MLP both gates finish in under half a minute at near-idle utilisation (the wall-clock is fixed startup, not method compute), so no meaningful ratio exists there and the cost question is decided entirely on the deep backbone. The ratios are the transferable claim; the absolute numbers are specific to this card, and because $\sim 15\text{--}30$ s of each run is startup, the training-only ratio is if anything slightly above $16\times$.

Dataset	Gate	wall [s]	util [%]	power [W]	energy [Wh]	mem [MB]
CIFAR-10 (ResNet-18)	Curriculum (scalar)	198 ± 7	49	259	13	4,203
	Asym. PER ($\delta=0.1$)	$3,157 \pm 11$	96	367	320	4,679
rot-MNIST (MLP)	Curriculum (scalar)	25 ± 1	< 1	—	0.4	749
	Asym. PER ($\delta=0.1$)	29 ± 1	8	—	0.5	769

Table C.10: Compute cost of one full two-task run (mean over five seeds, $\pm$ sd on wall-clock; single RTX PRO 6000; per-step evaluation disabled): wall-clock, mean GPU utilisation and power, integrated GPU energy, and peak GPU memory. On the ResNet-18 the directional gate costs $\sim 16\times$ the wall-clock and $\sim 24\times$ the energy of the scalar gate at similar peak memory. The MLP rows are dominated by fixed startup on a near-idle GPU, so their power draw is idle draw (omitted) and no ratio is meaningful.

C.9 Long-Sequence rot-MNIST (L-series)

Condition	ACC	FORG	min-ACC	WF ₁₀₀	WP ₁₀₀	WC-ACC	Depth	Area end
<i>Momentum 0.0</i>								
L1 vanilla ER	0.861 ± 0.004	0.060 ± 0.004	0.670 ± 0.055	0.207 ± 0.053	0.443 ± 0.011	0.722 ± 0.043	21.0 ± 5.8	3.14 ± 0.77
L2 curriculum $\lambda_{\min}=0.20$	0.860 ± 0.002	0.031 ± 0.005	0.828 ± 0.007	0.073 ± 0.005	0.400 ± 0.003	0.840 ± 0.005	4.1 ± 0.8	1.62 ± 0.42
L3 asym. PER $\delta=0.1$	0.870 ± 0.003	0.053 ± 0.006	0.843 ± 0.008	0.038 ± 0.007	0.381 ± 0.003	0.862 ± 0.007	4.7 ± 1.2	0.38 ± 0.24
<i>Momentum 0.9</i>								
L1 vanilla ER	0.889 ± 0.007	0.096 ± 0.008	0.845 ± 0.007	0.065 ± 0.008	0.317 ± 0.006	0.870 ± 0.006	8.9 ± 1.7	1.68 ± 0.69
L2 curriculum $\lambda_{\min}=0.20$	0.897 ± 0.002	0.071 ± 0.003	0.871 ± 0.005	0.041 ± 0.004	0.364 ± 0.005	0.888 ± 0.004	5.4 ± 0.6	0.60 ± 0.14
L3 asym. PER $\delta=0.1$	0.884 ± 0.004	0.104 ± 0.004	0.849 ± 0.006	0.054 ± 0.004	0.325 ± 0.005	0.873 ± 0.005	7.5 ± 1.5	0.80 ± 0.63

Table C.11: Five-task rot-MNIST (rotations $[0, 30, 60, 90, 120]^\circ$, four consecutive transitions, five seeds), the L-series: both gates in their practical configurations against vanilla ER, at both momentum settings. Depth is the worst single-boundary drop and area the end-referenced transient, aggregated as in the two-task blocks. At momentum 0 both gates cut the depth from 21 pp to ≈ 4 pp; under momentum the curriculum keeps the best cumulative retention while the re-inflated directional gate falls slightly behind vanilla on FORG (0.104 vs. 0.096, all five seeds).




Figure C.3: Five-task rot-MNIST (rotations $[0, 30, 60, 90, 120]^\circ$, four consecutive switches, five seeds with $\pm$ std bands), the L-series: first-task T_0 accuracy across the whole sequence for vanilla ER (red), the practical curriculum $\lambda_{\min}=0.20$ (blue) and asymmetric PER $\delta=0.1$ (aqua); dashed lines mark the four switches, and the two panels are the two momentum settings. (a) Momentum 0: both gates hold T_0 almost flat through every boundary, asymmetric PER the smoothest, while vanilla ER drops sharply at each switch. (b) Momentum 0.9: PER’s boundary transients return as momentum re-inflates its per-step filter (Section 6.3); the curriculum keeps the boundaries smoothest and holds the highest plateau (Table C.11).

D Derivation of the Trajectory Bound

This appendix derives the two results used in Section 3.1: the replay loss is monotone along the reference path (Proposition 1), and a gradient-flow iterate under a linear ramp tracks that path with $O(1/N)$ lag (Proposition 2), which together give the transient bound of Equation (3.2).

Setup. For $\lambda \in [0, 1]$ define the curriculum objective

$$L_\lambda(\theta) = L_{\text{replay}}(\theta) + \lambda L_{\text{current}}(\theta), \quad H_\lambda(\theta) = H_{\text{replay}}(\theta) + \lambda H_{\text{current}}(\theta), \quad (\text{D.1})$$

with $H_{\text{replay}} = \nabla^2 L_{\text{replay}}$ and $H_{\text{current}} = \nabla^2 L_{\text{current}}$. We make two assumptions.

- (A1) L_{replay} and L_{current} are C^2 in a neighbourhood of the path considered below.
- (A2) There is a C^1 branch of critical points $\lambda \mapsto \Theta^*(\lambda)$ of L_λ with $\Theta^*(0) = \theta_0^*$ and $H_\lambda(\Theta^*(\lambda)) \succ 0$ for all $\lambda \in [0, 1]$.

Assumption (A2) asks for a tracked branch of strict local minima, not global convexity. Near $\lambda = 0$ it holds automatically: θ_0^* is a non-degenerate local minimum of L_{replay} , so $H_0(\theta_0^*) \succ 0$ and the implicit function theorem yields a unique C^1 branch locally. Its persistence is controlled by Weyl’s inequality, $\mu_{\min}(H_\lambda) \geq \mu_{\min}(H_{\text{replay}}) + \lambda \mu_{\min}(H_{\text{current}})$, with $\mu_{\min}$ the smallest eigenvalue: positive-definiteness survives until the new task’s negative curvature outweighs the old task’s, and holds for all λ whenever $H_{\text{current}} \geq 0$ along the branch. On a deep non-convex backbone (A2) is a local modelling assumption, which is why Section 3.1 presents the bound as motivation rather than a guarantee.

Proposition 1 (Monotonicity of the replay loss). *Under (A1)–(A2),*

$$\frac{d}{d\lambda} L_{\text{replay}}(\Theta^*(\lambda)) = \lambda \nabla L_{\text{current}}^\top H_\lambda^{-1} \nabla L_{\text{current}} \geq 0 \quad \text{for all } \lambda \in [0, 1],$$

where $\nabla L_{\text{current}}$ and H_λ are evaluated at $\Theta^*(\lambda)$. In particular $L_{\text{replay}}(\Theta^*(\lambda)) \leq L_{\text{replay}}(\theta_{\text{joint}}^*)$ for all λ , with $\theta_{\text{joint}}^* = \Theta^*(1)$.

Proof. The branch satisfies the first-order condition

$$\nabla L_{\text{replay}}(\Theta^*(\lambda)) + \lambda \nabla L_{\text{current}}(\Theta^*(\lambda)) = 0. \quad (\text{D.2})$$

Differentiating (D.2) in λ and collecting terms gives $H_\lambda \frac{d\Theta^*}{d\lambda} + \nabla L_{\text{current}} = 0$, hence

$$\frac{d\Theta^*}{d\lambda} = -H_\lambda^{-1} \nabla L_{\text{current}}. \quad (\text{D.3})$$

By the chain rule and (D.2),

$$\frac{d}{d\lambda} L_{\text{replay}}(\Theta^*(\lambda)) = \nabla L_{\text{replay}}^\top \frac{d\Theta^*}{d\lambda} = (-\lambda \nabla L_{\text{current}})^\top (-H_\lambda^{-1} \nabla L_{\text{current}}) = \lambda \nabla L_{\text{current}}^\top H_\lambda^{-1} \nabla L_{\text{current}},$$

which is non-negative because $\lambda \geq 0$ and $H_\lambda^{-1} \succ 0$ by (A2). Integrating from λ to 1 gives the final statement. $\square$

Proposition 2 ($O(1/N)$ tracking lag). *Let $\theta(t)$ follow the gradient flow $\dot{\theta} = -\nabla L_{\lambda(t)}(\theta)$ under the linear ramp $\lambda(t) = t/N$, with $\theta(0) = \theta_0^*$, and let $e(t) = \theta(t) - \Theta^*(\lambda(t))$ denote the tracking error. Suppose, in addition to (A1)–(A2), that along the branch $\mu_{\min}(H_\lambda) \geq \mu > 0$ and $\|\nabla L_{\text{current}}\| \leq g$. Then, to first order in e , the error settles at*

$$\|e\| \leq \frac{1}{N} \|H_\lambda^{-1}\| \left\| \frac{d\Theta^*}{d\lambda} \right\| \leq \frac{g}{\mu^2 N} = O\left(\frac{1}{N}\right).$$

Proof. Differentiating the error and inserting the flow and the chain rule for the moving target,

$$\dot{e} = -\nabla L_{\lambda(t)}(\theta(t)) - \frac{d\Theta^*}{d\lambda} \dot{\lambda}, \quad \dot{\lambda} = \frac{1}{N}.$$

Linearising the gradient about the minimiser, $\nabla L_\lambda(\theta) = H_\lambda e + O(\|e\|^2)$ since $\nabla L_\lambda(\Theta^*) = 0$, yields the error dynamics

$$\dot{e} \approx -H_\lambda e - \frac{1}{N} \frac{d\Theta^*}{d\lambda}.$$

Because $H_\lambda \succeq \mu I$, these dynamics contract toward the quasi-static equilibrium $e^\dagger = -\frac{1}{N} H_\lambda^{-1} \frac{d\Theta^*}{d\lambda}$; bounding $\|H_\lambda^{-1}\| \leq 1/\mu$ and, by (D.3), $\|d\Theta^*/d\lambda\| \leq g/\mu$ gives the stated bound. The argument is a first-order, continuous-time approximation: it captures how the lag scales with the ramp length N , not an exact discrete-time constant. $\square$

The transient bound. Expanding the replay loss about the reference point, $L_{\text{replay}}(\theta(t)) = L_{\text{replay}}(\Theta^*) + \nabla L_{\text{replay}}(\Theta^*)^\top e + O(\|e\|^2)$, Proposition 1 bounds the first term by $L_{\text{replay}}(\theta_{\text{joint}}^*)$ and Proposition 2 bounds the correction by $O(1/N)$ (the gradient norm $\|\nabla L_{\text{replay}}(\Theta^*)\| = \lambda \|\nabla L_{\text{current}}\| \leq g$ is bounded along the branch). Hence

$$\max_t L_{\text{replay}}(\theta(t)) \leq L_{\text{replay}}(\theta_{\text{joint}}^*) + O\left(\frac{1}{N}\right): \quad (\text{D.4})$$

the permanent price is the first term, and the transient contribution of the tracking lag shrinks as the ramp lengthens, which is Equation (3.2).